



TOR VERGATA
UNIVERSITÀ DEGLI STUDI DI ROMA

Macroarea di Ingegneria

A QoS-Aware Broker for Multi-Provider Serverless Applications

Andrea Graziani

m. 0273395

Supervisor: Valeria Cardellini
Tutor: Gabriele Russo Russo

May 23, 2022

My thesis is focused on “**serverless computing**”.

It is a development paradigm according to which:

- The provider takes care of all aspect of server management.
- Cloud application are abstracted as a group of so-called **serverless functions**, which are computation units implementing a business functionality.

A serverless function is executed inside a containerized environment: the so-called **function instance**.

The FaaS platform **automatically** scales the number of function instances.

FaaS platforms impose a **limit** on the number of function instance runnable at the **same time** called **concurrency limit**.

A delay is observed when a new function instance is started by the provider: this event is called **cold start**.

To invoke a **serverless function**, users have to specify a so-called **serverless function configuration**.

- Generally, the **amount of memory allocated** to a serverless function.

Configuration parameters **significantly** affect the **cost** and **response time** of serverless functions.

- Lack of support for application whose functions are hosted on **multiple providers**.
- Lack of support for serverless function implementations abstraction, that is, for the so-called **concrete functions**.
- Fulfillment of **non-functional requirements** concerning the **quality of service** (QoS) levels that should be guaranteed for **multi-provider serverless applications**.
- Fulfillment of **functional requirements** concerning the **orchestration** of multi-provider applications.

Goal # 1

To guarantee the **satisfaction of QoS levels** for multi-provider serverless applications.

It was necessary to develop:

- 1 An **analytical model** to evaluate aforementioned class of applications.
- 2 A **methodological way** to find the “best” configuration to satisfy QoS constraints.


Goal # 2


The **orchestration** for multi-provider serverless applications.


To achieve it, I had to build:

- 1 A **software framework**.
- 2 An extension to an already existent **representation scheme** to define a serverless application workflow.
 - Based on an existing language called **abstract function choreography language** (AFCL).

Serverless applications are abstracted to a **weakly connected weighted directed graph**.







AWS
Lambda

f_x Concrete Functions 1

f_x Concrete Functions 2

f_x Concrete Functions 3

...





APACHE
OpenWhisk

f_x Concrete Functions 1


f_x Concrete Functions 2

...




AWS Lambda

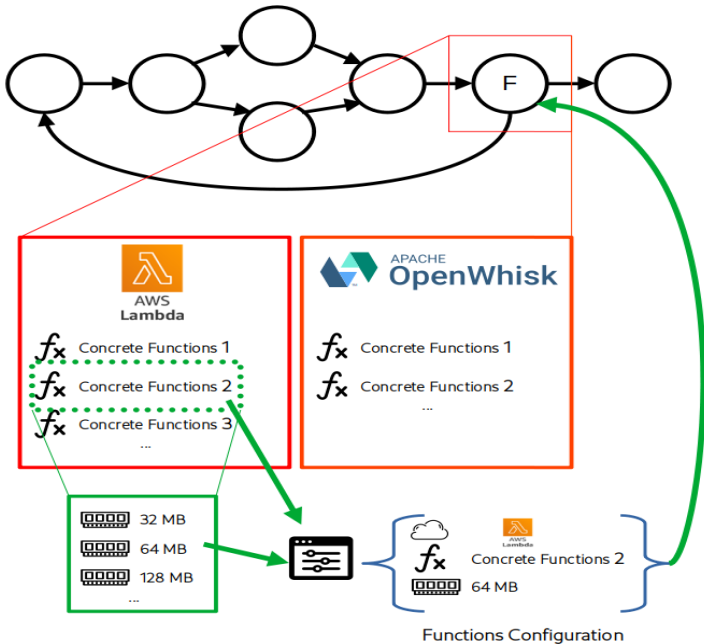
f_x Concrete Functions 1
 f_x Concrete Functions 2
 f_x Concrete Functions 3
 ...



APACHE OpenWhisk

f_x Concrete Functions 1
 f_x Concrete Functions 2
 ...

 32 MB
 64 MB
 128 MB
 ...



- Estimations of **average response time** and **cost** under **any** possible configuration of all concrete functions is done using **exponential moving average** approach.
- An estimation of the **probability** according to which a request follows a cold start is required.
 - This is done using **Erlang-B** formula by modeling FaaS platform providers by **sets** of $M/G/K(t)/K(t)$ queueing systems.
 - $K(t)$: the number of function instances at time t .
- Performance estimations of the application depend on its workflow properties.

- To achieve our goal consisting in finding the best configuration to guarantee QoS constraints, we have to solve an **optimization problem**.
 - It is based on **multi-dimensional multi-choice knapsack problem formulation** (MMKP).

$$\max \sum_{i=1}^{|\mathcal{F}_{\mathcal{E}}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot p_{\phi_{i_j}}(t) \quad (5.10)$$

$$\text{subject to } \sum_{i=1}^{|\mathcal{F}_{\mathcal{E}}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot c_{\phi_{i_j}}(t) \leq C \quad (5.11)$$

$$\begin{aligned} & \sum_{\phi_i \in \mathcal{F}_{\mathcal{E}}(C) \setminus \mathcal{F}_{\mathcal{E}}(\tilde{\mathcal{P}})} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot rt_{\phi_{i_j}}(t) + \\ & + \sum_{\phi_h \in \delta_C} \sum_{j=1}^{|\mathbf{F}_{\phi_h} \times \mathbb{N}|} y_{\phi_{h_j}} \cdot rt_{\phi_{h_j}}(t) \leq RT \end{aligned} \quad \forall \delta_C \in \Delta_C \quad (5.12)$$

$$\sum_{i=1}^{|\mathcal{F}_{\mathcal{E}}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot a_{(\phi_{i_j}, \omega_P^{(l)})} \leq l - R(\mathbf{Q}_{\omega_P^{(l)}}, t) \quad \forall \omega_P^{(l)} \in \tilde{\mathbf{S}}_C \quad (5.13)$$

$$\sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} = 1 \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_{\mathcal{E}}|] \quad (5.14)$$

$$y_{\phi_{i_j}} \in \{0, 1\} \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_{\mathcal{E}}|] \\ \forall j \in \mathbb{N} \cap [1, |\mathbf{F}_{\phi_i} \times \mathbb{N}|]$$

$$\max \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot p_{\phi_{ij}}(t) \quad (5.10)$$

$$\text{subject to } \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot c_{\phi_{ij}}(t) \leq C \quad (5.11)$$

$$\begin{aligned} & \sum_{\phi_i \in \mathcal{F}_\mathcal{E}(\mathcal{C}) \setminus \mathcal{F}_\mathcal{E}(\tilde{\mathcal{P}})} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot rt_{\phi_{ij}}(t) + \text{Cost Constraint} \\ & + \sum_{\phi_h \in \delta_C} \sum_{j=1}^{|\mathbf{F}_{\phi_h} \times \mathbb{N}|} y_{\phi_{hj}} \cdot rt_{\phi_{hj}}(t) \leq RT \end{aligned} \quad \forall \delta_C \in \Delta_C \quad (5.12)$$

$$\sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot a_{(\phi_{ij}, \omega_P^{(l)})} \leq l - R(\mathbf{Q}_{\omega_P^{(l)}}, t) \quad \forall \omega_P^{(l)} \in \tilde{\mathbf{S}}_C \quad (5.13)$$

$$\sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} = 1 \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \quad (5.14)$$

$$y_{\phi_{ij}} \in \{0, 1\} \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \\ \forall j \in \mathbb{N} \cap [1, |\mathbf{F}_{\phi_i} \times \mathbb{N}|]$$

$$\max \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot p_{\phi_{i_j}}(t) \quad (5.10)$$

$$\text{subject to } \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot c_{\phi_{i_j}}(t) \leq C \quad (5.11)$$

$$\begin{aligned} & \sum_{\phi_i \in \mathcal{F}_\mathcal{E}(\mathcal{C}) \setminus \mathcal{F}_\mathcal{E}(\tilde{\mathcal{P}})} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot rt_{\phi_{i_j}}(t) + \\ & + \sum_{\phi_h \in \delta_\mathcal{C}} \sum_{j=1}^{|\mathbf{F}_{\phi_h} \times \mathbb{N}|} y_{\phi_{h_j}} \cdot rt_{\phi_{h_j}}(t) \leq RT \end{aligned} \quad \begin{array}{l} \text{Response Time Constraint} \\ \forall \delta_\mathcal{C} \in \Delta_\mathcal{C} \end{array} \quad (5.12)$$

$$\sum_{i=1}^{|\mathcal{F}_\mathcal{E}(\mathcal{C})|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} \cdot a_{(\phi_{i_j}, \omega_P^{(l)})} \leq l - R(\mathbf{Q}_{\omega_P^{(l)}}, t) \quad \forall \omega_P^{(l)} \in \tilde{\mathbf{S}}_\mathcal{C} \quad (5.13)$$

$$\sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{i_j}} = 1 \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \quad (5.14)$$

$$y_{\phi_{i_j}} \in \{0, 1\} \quad \begin{array}{l} \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \\ \forall j \in \mathbb{N} \cap [1, |\mathbf{F}_{\phi_i} \times \mathbb{N}|] \end{array}$$

$$\max \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot p_{\phi_{ij}}(t) \quad (5.10)$$

$$\text{subject to } \sum_{i=1}^{|\mathcal{F}_\mathcal{E}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot c_{\phi_{ij}}(t) \leq C \quad (5.11)$$

$$\begin{aligned} & \sum_{\phi_i \in \mathcal{F}_\mathcal{E}(C) \setminus \mathcal{F}_\mathcal{E}(\tilde{\mathcal{P}})} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot rt_{\phi_{ij}}(t) + \\ & + \sum_{\phi_h \in \delta_C} \sum_{j=1}^{|\mathbf{F}_{\phi_h} \times \mathbb{N}|} y_{\phi_{hj}} \cdot rt_{\phi_{hj}}(t) \leq RT \end{aligned} \quad \forall \delta_C \in \Delta_C \quad (5.12)$$

$$\sum_{i=1}^{|\mathcal{F}_\mathcal{E}(C)|} \sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} \cdot a_{(\phi_{ij}, \omega_P^{(l)})} \leq l - R(\mathbf{Q}_{\omega_P^{(l)}}, t) \quad \forall \omega_P^{(l)} \in \tilde{\mathbf{S}}_C \quad (5.12)$$

$$\sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} y_{\phi_{ij}} = 1 \quad \text{Capacity Constraint} \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \quad (5.14)$$

$$y_{\phi_{ij}} \in \{0, 1\} \quad \forall i \in \mathbb{N} \cap [1, |\mathcal{F}_\mathcal{E}|] \quad \forall j \in \mathbb{N} \cap [1, |\mathbf{F}_{\phi_i} \times \mathbb{N}|]$$

I develop a custom heuristic algorithm based on **ant colony optimization** (ACO).

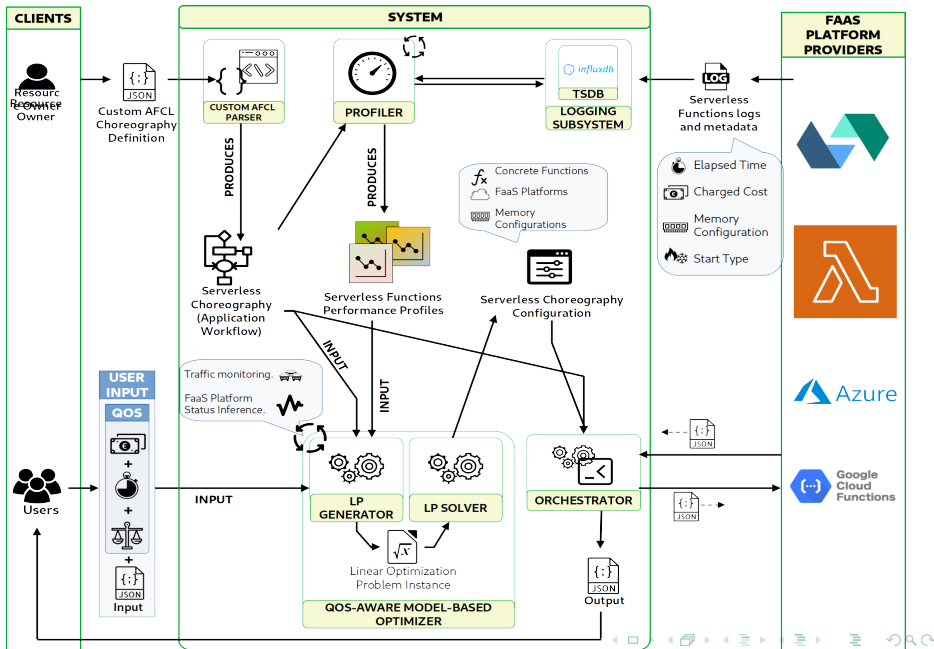
- It is based on a set of computational agents, called **artificial ants**, which **iteratively** construct a solution.
- At each iteration, each agent moves from a solution to another, applying a series of stochastic **local** decisions whose policy is based on following parameters:
 - **Attractiveness.**
 - **Pheromone trail.**

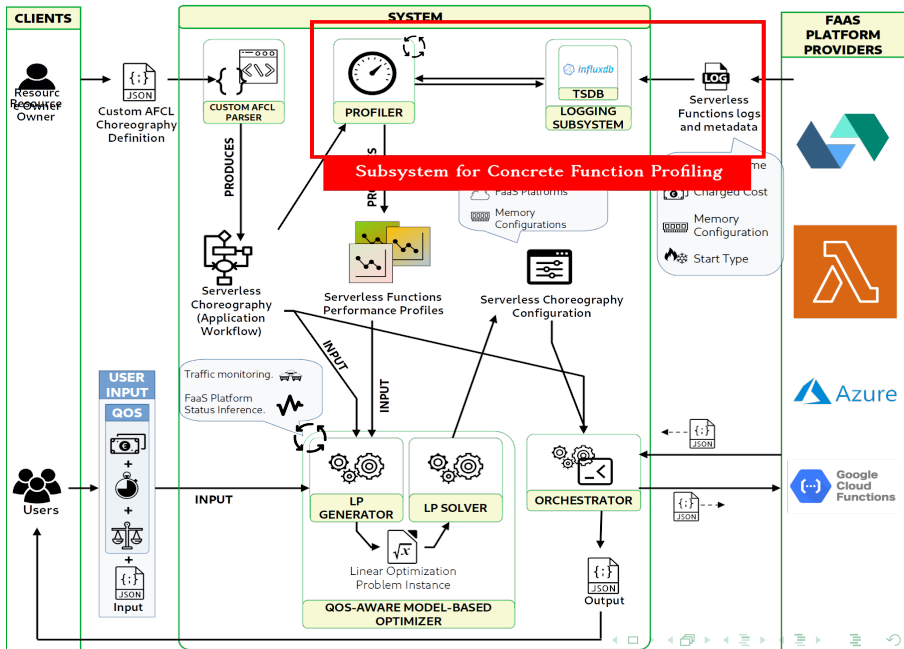
Pheromone trails are updated during **every iteration**.

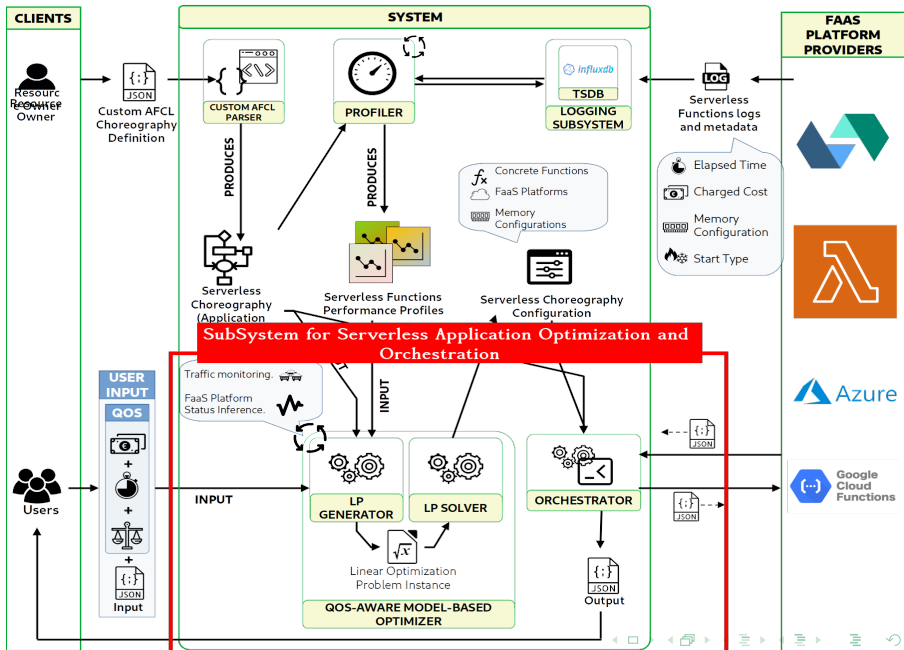
Pheromone trails are used to decide which solutions should be preferred during **next iterations**.

Main features of our software framework:

- **Client-server architecture.**
- **Cloud-native** application.
- Includes a set of **adapters** to interact with following FaaS providers:
 - AWS Lambda.
 - Apache OpenWhisk.
- **REST** architectural style.









We validate our model through several experiments using an image-processing serverless application.

- Firstly, we check the respect of user specified QoS constraints in a static way thorough several **sequential invocations**.
- Then, we test the model in a dynamically way thorough **several concurrent and parallel invocations**.
 - That experiment aimed to **run out of capacity** on one FaaS provider and so force our prototype to schedule the concrete function on the other.


- Application configuration produced by our system **before** Provider 1 runs out of capacity...




- ...and **after**, when Provider 1 runs out of capacity.



- We compared **execution time** of our heuristic algorithm with that of optimal algorithm through several experiments.



- According to my results, **accuracy** is above **98%**.



- 1 I presented an analytical model to evaluate the performance of multi-provider serverless application.
- 2 I defined an optimization problem formulation to address the problem of finding a suitable configuration in order to meet user defined QoS constraints.
- 3 I developed a heuristic algorithm to rapidly solve it.
- 4 I validate proposed solution through test and experimental evaluations using my prototype.



Thanks for your attention!
Questions?