

References

- [1] Nabeel Akhtar, Ali Raza, Vatche Ishakian, and Ibrahim Matta. Cose: Configuring serverless functions using statistical learning. In *IEEE INFOCOM 2020 - IEEE Conference on Computer Communications*, pages 129–138, July 2020.
- [2] Nima Mahmoudi and Hamzeh Khazaei. Performance modeling of serverless computing platforms. *IEEE Transactions on Cloud Computing*, pages 1–1, 2020.
- [3] Xiongfei Weng, Hongliang Yu, Guangyu Shi, Jian Chen, Xu Wang, Jing Sun, and Weimin Zheng. Understanding locality-awareness in peer-to-peer systems. In *2008 International Conference on Parallel Processing - Workshops*, pages 59–66, 2008.

Contents

1	Serverless Computing Paradigm	4
1.1	Serverless Function	4
1.2	Serverless Application	4
1.3	fdsfsd	5
1.4	title	5
1.5	Concurrency level	6
1.6	Start/Cold start	6
1.6.1	Cold start probability	7
2	Orchestrator Model	7
2.1	Resource owner	7
3	Serverless Choreography	8
3.1	Preliminary Definitions	8
3.2	Definition	9
3.2.1	Abstract Serverless Function	10
3.2.2	Executability condition	11
3.2.3	Serverless Sub-choreography	12
3.2.4	Basic Serverless Choreography	12
4	Serverless Choreography Configuration	13
4.1	Executable Function Configuration	13
4.2	Serverless Choreography Configuration	13
5	The problem of cold starts	14
5.1	Metrics	14
6	Serverless Choreography Structures	14
6.0.1	Parallel	15
6.0.2	Branch/Switch	15
6.0.3	Loop	15

7 Optimization	16
7.1 Performance Modeling	16
8 Super-peer Node	17
9 Peer Node	17
10 Resources	18
11 System's resources and actors	18
11.1 Server-less function swarms	18
11.1.1 Server-less function sub-swarms	18
12 Function choreography scheduler	19
12.1 Schedulability condition	19
12.2 The Δ_{X_R} -Scheduler	19
12.2.1 Virtual function instance	20
12.2.2 Proprieties and constrains	20
12.3 The FC_R -Scheduler	20
13 FC_R-Active Peer Node	21
14 Architecture overview	21
14.1 Overlay network	21
14.1.1 Locality-awareness property	21
14.2 Fault Tolerance	22
14.2.1 Availability	22
14.2.2 Replication	22
14.3 FC_R -Request	23
15 Queuing system of a Δ_{X_R}-scheduler	24
15.1 Server-less function preemption	24
15.2 Queuing system design	24
15.3 The virtual function instances allocation problem	25
15.3.1 Reactive scaling policy	25
15.3.2 Server-less function scheduling for QoS "Minimum Response Time"	26

fdsfsdfsdfsdsd f sd fsd f sdfsdsdfsdsdterterterterterterterterterterterter

$$\begin{aligned}
\min \quad & \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} g(x)^{f_k^i} Y_x^{f_k^i} \\
& \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} T(x)^{f_k^i} Y_x^{f_k^i} \leq D_{\mathbf{F}} \\
& \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} Y_x^{f_k^i} = 1
\end{aligned} \tag{1}$$

$$Y_x^{f_k^i} \in \{0, 1\}$$

$$\begin{aligned}
\min \quad & \sum_{\pi \in \Pi} \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} g(x)^{f_k^i} Y_{(\pi, x)}^{f_k^i} \\
& \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} T(x)^{f_k^i} Y_x^{f_k^i} \leq D_{\mathbf{F}} \\
& \sum_{\pi \in \Pi} \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} Y_{(\pi, x)}^{f_k^i} = 1
\end{aligned} \tag{2}$$

$$Y_{(\pi, x)}^{f_k^i} \in \{0, 1\}$$

1 Serverless Computing Paradigm

In serverless computing platforms, computation is done by so-called *function instances* which are completely managed by the serverless computing platform provider and act as tiny servers that can be invoked based on events forwarded by end-users [2].

Serverless computing platforms handle almost every aspect of the system administration tasks needed to deploy a workload on the cloud, providing a new simplified programming model according to which developers can focus on the business aspects of their applications only [1].

Moreover, the paradigm lowers the cost of deploying applications too, adopting a so-called “*pay as you go*” pricing model, by charging for execution time rather than for allocated resources [1].

1.1 Serverless Function

A *serverless function* represents a stateless, event-driven, self-contained unit of computation implementing a business functionality.

Although a serverless function generally represents a unit of executable code, submitted to FaaS platforms by developers using one or a combination of the programming languages supported by FaaS providers, a serverless function can be any cloud services eventually necessary to business logic, like cloud storage, message queue services, pub/sub messaging service etc.

When it represents executable code, developers can specify several configuration parameters, like timeout, memory size or CPU power [1].

A serverless function can be triggered through events or HTTP requests following which the FaaS provider executes it in a containerized environment, like container, virtual machine or even processes, using the specified configuration.

1.2 Serverless Application

In serverless paradigm, the most basic scenario is to invoke a single function through a HTTP request; however, if you want to build more complex applications, constructing so-called *serverless application*, serverless functions must be connected and appropriately coordinated.

Formally, we define a serverless application as a stateless and event-driven software system made up of a serverless functions set hosted on one or more FaaS platforms and combined together by a so-called *coordinator* (or *orchestrator*).

Generally, the coordinator is a broker needed to implement the business logic of any application: it chains together serverless function, handles events among functions and triggers them in the correct order according to the business logic defined by developers.

Most public cloud providers for serverless computing have introduced platforms to define and coordinate serverless function in order to build serverless application, like AWS Step Functions which combines multiple Lambda functions and other serverless services offered by AWS into responsive serverless applications

Although FaaS platforms continuously advance the support for serverless applications, existing solutions are dominated by a few large-scale providers resulting in mostly non-portable FaaS solutions

and provider lock-in

occurs when transitioning data, products, or services to another vendor’s platform is difficult and costly, making customers more dependent (locked-in) on a single cloud storage solution.

FC languages and runtime systems are still in their infancy with numerous drawbacks. Current commercial and open source offerings of FC systems are dominated by a few large-scale providers who offer their own platforms,

As FaaS platforms take over operational responsibilities, besides uploading the source code of functions, users have limited control over resources on FaaS platforms. Taking AWS Lambda as an example, the amount of allocated memory during execution and the concurrency level are only options for tuning the performance of functions. The amount of allocated memory is between 128 MB and 3,008 MB in 64MB increments [13]. Previous researches have proven that computational power and network throughput are in proportion to the amount of allocated memory, and disk performance also increases with larger memory size due to less contention [14], [15]. By reserving and provisioning more instances to host functions, high concurrency level can decrease fluctuations in the function performance incurred by cold starts (container initialization provisioning delay if no warm instance is available) and reduce the number of throttles under very heavy request loads

1.3 fdfsfd

•

1.4 title

Serverless computing has given a much-needed agility to developers, abstracted away the management and maintenance of physical resources, and provided them with a relatively small set of configuration parameters: memory and CPU. While relatively simpler, configuring the “best” values for these parameters while minimizing cost and meeting performance and delay constraints poses a new set of challenges. This is due to several factors that can significantly affect the running time of serverless functions.

the run-time of these serverless functions decreases with the increase of memory size allocated to the function. However, the marginal improvement in the run-time decreases as the memory increases.

This behavior is because the pricing model as exposed by the cloud providers is tightly coupled with the amount of resources specified to execute the serverless function (c.f. Figure 1a), and the dependency between memory and CPU resource allocation – AWS Lambda allocates CPU power

a model for each executable belonging to a given choreography.

This helps developers decide if a developed workload would comply with their QoS agreements, and if not, how much performance improvement they would need to do so. The performance improvement decided could be achieved either by improving the design, quality of code, or by simply resizing the resource allocated to each instance, which is usually set by changing the allocated memory.

A cold start

Cold/Warm start: as defined in previous work [3], [8], [10], we refer to cold start request when the request goes through the process of launching a new function instance. For the platform, this could include launching a new virtual machine, deploying a new function, or creating a new instance on an existing virtual machine, which introduces an overhead to the response time experienced by users. In case the platform has an instance in the idle state when a new request arrives, it will reuse the existing function instance instead of spinning up a new one. This is commonly known as a warm start request. Cold starts could be orders of magnitude longer than warm starts for some applications. Thus, too many cold starts could impact the application’s responsiveness and user experience [3].

his equation gives the probability of a request being rejected (blocked) by the warm pool, assuming there are m warm servers. If m is less than the maximum concurrency level, the request blocked by the warm pool causes a cold start. If the warm pool has reached the maximum concurrency level, any request rejected by the warm pool will be rejected by the platform.

1.5 Concurrency level

Any public serverless computing platform imposes a limitation on the number of serverless function instance runnable at the same time by the corresponding platform provider; this limit is known as *concurrency level*.

For example, in 2021, AWS Lambda doesn’t allow more than 1000 serverless function instances in running state at the same time.

1.6 Start/Cold start

When a request arrives on a FaaS platform,

Cold start If the warm pool is busy, that is there are no serverless function instances in idle state able to serve an newly incoming request, latter will trigger the launching of a new function instance, which will be added to the warm pool. This event is called *cold start*.

From the FaaS provider point of view, this operation requires to start either a new virtual machine or a new container; in any case, regardless of the method adopted, a cold start introduces a very important overhead to the response time experienced by users.

Warm start When there is at least one serverless function instance in idle state, the FaaS platform reuses it to serve the incoming request without launching a new one. This is called *warm start*.

A very important aspect is that cold starts could be orders of magnitude longer than warm starts for some applications; therefore, too many cold starts could impact the application’s responsiveness and user experience.

To be more precise, let $k \in \mathbb{N}$ and $C_{\max, P}$ the maximum concurrency level imposed by FaaS provider P , when a new request arrives on the platform, one of the following events can occur:

- If $k < C_{\max, P}$, that is the number of function instances into the warm pool is less than the maximum concurrency level, the request will be blocked by the system causing a cold start.

- If $k = C_{\max, P}$, that is the warm pool size reaches the maximum concurrency level, any further request will be rejected by the platform

1.6.1 Cold start probability

As previously said, the rejection of a request by the warm pool triggers a cold start, adding subsequently a new function instance to the warm pool in order to handle received requests.

To build our framework, we must to know the probability of a request being rejected by the warm pool; in other words, we must to compute the *cold start probability*.

Formally, let $k \in N$, the cold start probability \mathbf{P}_C is the probability that an arrival finds all k function instances of the warm pool busy. Using Erlang-B formula, aforementioned probability can be calculated as follows:

$$\mathbf{P}_C = \frac{\frac{\rho^k}{k!}}{\sum_{j=0}^k \frac{\rho^j}{j!}} \quad (3)$$

2 Orchestrator Model

2.1 Resource owner

A *resource owner*, henceforward denoted with R , represents an entity capable of *creating*, *modifying* and *authorizing* access to several resources of our system.

3 Serverless Choreography

In this section we will introduce the notion of *serverless choreography*, a very important *resource* adopted to model and implement both serverless functions and serverless applications.

Informally, the notion of serverless choreography has been derived from that of a control-flow graph; as known, the latter describes, using graph notation, all paths that might be traversed through a serverless application during its execution. Similarly, a serverless choreography describes calling relationships between serverless functions which can be combined using several types of control-flow connectors.

3.1 Preliminary Definitions

Before to define what we mean for serverless choreography, is necessary to introduce some very useful notations.

Let $n \in \mathbb{N}$ and $G = (\Phi, E)$ a directed graph, where:

- Φ is a finite set of vertices, such that $|\Phi| = n$;
- $E \subseteq \Phi \times \Phi$ is a finite set of ordered pairs of vertices $e_{ij} = (\phi_i, \phi_j)$, where $\phi_i \in \Phi$ to $\phi_j \in \Phi$ for any $i, j \in \mathbb{N} \cap [1, n]$. Any ordered pair of vertices is also called directed edge;

Then, we will adopt following notations:

- A *path* of G is defined as a finite sequence of distinct vertices and edges. We will denote a path by π which formally can be represented as follows:

$$\pi = \phi_1 e_1 \phi_2 e_2 \dots e_{n-2} \phi_{n-1} e_{n-1} \phi_n \quad (4)$$

where:

- $\phi_i \in \Phi$, for all $i \in \mathbb{N} \cap [1, n]$
- $e_i = (\phi_i, \phi_{i+1}) \in E$, for all $i \in \mathbb{N} \cap [1, n-1]$
- Let $\phi_i, \phi_j \in \Phi$ for any $i, j \in \mathbb{N} \cap [1, n]$, the set denoted by $\Pi(\phi_i, \phi_j)$ identifies all possible paths starting from vertex ϕ_i and ending at vertex ϕ_j .
- For any $u \in \mathbb{N} \cap [1, n]$, the set $out(\phi_u)$ ($in(\phi_u)$) denotes all edges starting (ending) from (to) vertex ϕ_u , while the set $succ(\phi_u)$ ($pred(\phi_u)$) includes all direct successor (predecessors) vertices of ϕ_u . Formally:

$$out(\phi_u) \stackrel{def}{=} \{(\phi_u, \phi) \in E, \quad \forall \phi \in \Phi\} \quad (5)$$

$$in(\phi_u) \stackrel{def}{=} \{(\phi, \phi_u) \in E, \quad \forall \phi \in \Phi\} \quad (6)$$

$$succ(\phi_u) \stackrel{def}{=} \{\phi \in \Phi \mid (\phi_u, \phi) \in out(\phi_u)\} \quad (7)$$

$$pred(\phi_u) \stackrel{def}{=} \{\phi \in \Phi \mid (\phi, \phi_u) \in in(\phi_u)\} \quad (8)$$

3.2 Definition

According to the serverless paradigm, the execution of a serverless application starts with a particular function, which we will call *entry point*, while any other serverless functions, belonging to the application, will be invoked subsequently according to specified business logic defined by developers; as we will see shortly, the latter can be naturally modeled by a weighted directed graph.

Clearly, the execution of a serverless application ends when the execution of the last function of the application ends; that ending function will be called as *end point*. We assume that the entry point is unique.

Let $n \in \mathbb{N} \setminus \{0\}$ and R a resource owner.

A *serverless choreography* owned by R , or simply *choreography* of R , is a resource represented by a weighted directed graph denoted as follows:

$$\mathcal{C}_R \stackrel{def}{=} (\Phi, E) \quad (9)$$

where:

- $|\Phi| = n$;
- Each vertex $\phi \in \Phi$ is called *abstract serverless function* and represents a computational unit; we will describe more in detail what we mean by abstract serverless function shortly;
- Let $i, j \in \mathbb{N} \cap [1, n]$ and $\phi_i, \phi_j \in \Phi$, any directed edge $e_{ij} = (\phi_i, \phi_j) \in E$ represents the calling relationship between two abstract serverless function which depends on the business logic defined by R .

In our context, any directed edge $(\phi_i, \phi_j) \in E$ states that the abstract serverless function ϕ_j can be called by ϕ_i ;

- Let $i, j \in \mathbb{N} \cap [1, n]$, the number $p_{ij} \in \mathbb{R} \cap [0, 1]$ is the weight assigned to the edge $(\phi_i, \phi_j) \in E$, where:
 - The number p_{ij} represents the so-called *transition probability* from ϕ_i to ϕ_j , that is the probability of invoking ϕ_j after finishing the execution of ϕ_i ;
 - We will use a function $P : \Phi \times \Phi \rightarrow [0, 1]$, called *transition probability function*, such that $P(\phi_i, \phi_j) = p_{ij}$. When $P(\phi_i, \phi_j) = 0$, it implies that the directed edge $(\phi_i, \phi_j) \notin E$, therefore ϕ_i cannot invoke ϕ_j ;
 - For any path $\pi = \phi_1 e_1 \dots e_{n-1} \phi_n$ of \mathcal{C}_R , we define *transition probability of the path π* the following quantity:

$$TPP(\pi) = \prod_{i=1}^{n-1} P(\phi_i, \phi_{i+1}) \quad (10)$$

- Φ contains only one serverless abstract function, denoted by ϕ_{entry} , acting as the entry point and at least one acting as the end point, which is denoted, instead, by ϕ_{end} . Formally, we define aforementioned vertices as follows:

$$\begin{aligned}
\phi \in \Phi & \text{ is the entry point of } \mathcal{S} & \Leftrightarrow & \text{in}(\phi) = \emptyset \\
\phi \in \Phi & \text{ is the end point of } \mathcal{S} & \Leftrightarrow & \text{out}(\phi) = \emptyset
\end{aligned} \tag{11}$$

Then, following conditions must be hold:

$$\exists! \phi \in \Phi \mid \text{in}(\phi) = \emptyset \tag{12}$$

$$\exists \phi \in \Phi \mid \text{out}(\phi) = \emptyset \tag{13}$$

Finally, we will also use the notation $\alpha(\mathcal{C}_R)$ to denote the entry point ϕ_{entry} of a choreography \mathcal{C}_R . Conversely, we will adopt the notation $\omega(\mathcal{C}_R)$ to denote the set of all end points of \mathcal{C}_R

- Must be hold the condition according to which the transition probabilities of all paths between the entry point and the end point of \mathcal{C}_R sum up to 1. Formally:

$$\sum_{\phi_{end} \in \omega(\mathcal{C}_R)} \left(\sum_{\pi \in \Pi(\phi_{entry}, \phi_{end})} TPP(\pi) \right) = 1 \tag{14}$$

In other words, above conditions guarantees that any execution starting from ϕ_{entry} will terminate.

- A function, denoted by $D : \Phi \times \Phi \rightarrow [0, \infty)$, represents a delay function according to which $D(\phi_i, \phi_j)$ identifies the delay from ϕ_i to ϕ_j due to network delay or orchestration task;

A choreography \mathcal{C}_R can be uniquely identified by an ordered pair (a, b) , where a is the name of the resource owner R , while b is the function choreography name.

Clearly, we say that choreography models a serverless function when $|\Phi| = 1$ and $|E| = 0$; conversely, it models a serverless function application when $|\Phi| > 1$ and $|E| > 0$.

From now, a choreography \mathcal{C}_R will be briefly denoted by \mathcal{C} when no confusion can arise about the resource owner R .

3.2.1 Abstract Serverless Function

Supposing that a choreography $\mathcal{C} = (\Phi, E)$ is given.

As said previously, any $\phi \in \Phi$ is called *abstract serverless function*, or simply *abstract function*, that is a *resource* representing a computational unit required by business logic provided by developers.

According to our model, there are two types of abstract functions implementations:

- ϕ is called *serverless executable functions*, or simply *executable function*, when ϕ contains executable code; therefore, in that case, we said that an executable function models a serverless function.

$\mathcal{F}_{\mathcal{C}}$ is defined as the set containing all executable function of \mathcal{C} and it is formally defined as follows:

$$\mathcal{F}_\varepsilon \stackrel{def}{=} \{\phi \in \Phi \mid \phi \text{ is a serverless executable function} \} \quad (15)$$

However, multiple different implementations of a given executable function can be provided by developers which, although they must be semantically and logically equivalent, may eventually expose different performance or cost behavior.

For any $\phi \in \mathcal{F}_\varepsilon$, we will use \mathbf{F}_ϕ notation to represent the so-called *implementation-set* of ϕ , that is the set containing any concrete implementation, denoted as f_ϕ , of ϕ .

Later, we will explain how our framework is able to pick, for all $\phi \in \mathcal{F}_\varepsilon$, exactly one $f_\phi \in \mathbf{F}_\phi$ whose properties allow us to meet user-specified QoS objective.

- ϕ is called *serverless orchestration functions*, or simply *orchestration functions*, when ϕ contains the so-called *orchestration code*.

According to our model, orchestration code represents the logic required to chain together any components of an application, handling events and triggering executable functions in the correct order according to the business logic. In other words, orchestration code is used to manage the control-flow of any application; later we will describe how it is possible.

$\mathcal{F}_\mathcal{O}$ is defined as the set containing all orchestration functions of \mathcal{C} and it is formally defined as follows:

$$\mathcal{F}_\mathcal{O} \stackrel{def}{=} \{\phi \in \Phi \mid \phi \text{ is a serverless orchestration function} \} \quad (16)$$

Clearly, based on above definitions, we can say:

$$\mathcal{F}_\varepsilon \cap \mathcal{F}_\mathcal{O} = \emptyset \quad (17)$$

$$\mathcal{F}_\varepsilon \cup \mathcal{F}_\mathcal{O} = \Phi \quad (18)$$

$$|\mathcal{F}_\varepsilon| + |\mathcal{F}_\mathcal{O}| = |\Phi| \quad (19)$$

Any abstract function ϕ is uniquely identified by an ordered pair (a, b) , where:

- a represents the identifier of the choreography \mathcal{C} ;
- b is the name of the abstract serverless function ϕ ;

3.2.2 Executability condition

Let $\mathcal{C} = (\Phi, E)$ a choreography.

Obviously, in order to effectively start the execution of a choreography, is required that, for each executable function $\phi \in \mathcal{F}_\varepsilon$, *at least one* concrete implementation f_ϕ exists.

Formally, we said that a choreography is *executable* when:

$$\mathcal{C} \text{ is executable} \Leftrightarrow |\mathbf{F}_\phi| \geq 1 \quad \forall \phi \in \mathcal{F}_\varepsilon \quad (20)$$

We will only deal with executable choreographies. Moreover, for all $\phi \in \mathcal{F}_\varepsilon$, we always assume that all $f_\phi \in \mathbf{F}_\phi$ are already deployed on one or more FaaS platform by developers.

3.2.3 Serverless Sub-choreography

Let $\mathcal{C} = (\Phi, E)$ a choreography.

The weighted directed sub-graph of \mathcal{C} defined as follows:

$$\mathcal{C}^* \stackrel{\text{def}}{=} (\Phi^*, E^*) \quad \text{where } \Phi^* \subseteq \Phi \wedge E^* \subseteq E \quad (21)$$

is called *serverless sub-choreography* of \mathcal{C} , or simply *sub-choreography* of \mathcal{C} , when the conditions 12, 13 and 14 are hold.

3.2.4 Basic Serverless Choreography

Suppose to have a choreography $\mathcal{C} = (\Phi, E)$ satisfying following conditions:

$$|\mathcal{F}_\emptyset| = 0 \quad (22)$$

$$|\omega(\mathcal{C})| = 1 \quad (23)$$

We said that any choreography \mathcal{C} , satisfying the conditions 22 and 23, represents a *basic serverless choreography*, or simply a *basic choreography*.

4 Serverless Choreography Configuration

Given a choreography $\mathcal{C} = (\Phi, E)$, the basic goal of our framework is to determine the so-called *serverless choreography configuration*, called also *choreography configuration* or, simply, *configuration*, which allows us to meet user-specified QoS objectives.

Informally, a configuration specifies which concrete implementation $f_\phi \in \mathbf{F}_\phi$, for all $\phi \in \mathcal{F}_\mathcal{E}$, will be effectively executed, including several parameters, like memory size or CPU power.

4.1 Executable Function Configuration

Before explaining formally what we meant by choreography configuration, we must firstly to explain the concept of executable function configuration.

Let $m \in \mathbb{N}$, $\phi \in \mathcal{F}_\mathcal{E}$ an executable function and \mathbf{F}_ϕ the corresponding implementation-set.

Formally, an *executable function configuration* for the executable function ϕ is a two-dimensional vector defined as follows:

$$x_\phi = (f_\phi, m) \in \mathbf{F}_\phi \times \mathbb{N} \quad (24)$$

where:

- f_ϕ represents a particular serverless function implementing the executable function ϕ .
- m represents the allocated memory size of f_ϕ .

At this point, we can define some useful functions:

- $C_{\mathbf{F}_\phi} : \mathbf{F}_\phi \times \mathbb{N} \rightarrow [0, \infty)$ is the *cost function* for any serverless functions belonging to the implementation-set \mathbf{F}_ϕ .

Precisely, for all $f_\phi \in \mathbf{F}_\phi$ and $m \in \mathbb{N}$, $C_{\mathbf{F}_\phi}(f_\phi, m)$ returns the *average cost* paid by developers to execute f_ϕ using an allocated memory size equal to m .

- $RT_{\mathbf{F}_\phi} : \mathbf{F}_\phi \times \mathbb{N} \rightarrow [0, \infty)$ is a *delay function* for any serverless functions belonging to the implementation-set \mathbf{F}_ϕ .

For all $f_\phi \in \mathbf{F}_\phi$ and $m \in \mathbb{N}$, $RT_{\mathbf{F}_\phi}(f_\phi, m)$ returns the *average response time* when f_ϕ is invoked with memory size equal to m ;

4.2 Serverless Choreography Configuration

Let $n, k \in \mathbb{N} \setminus \{0\}$ and a choreography $\mathcal{C} = (\Phi, E)$ such that $|\Phi| = n$ and $|\mathcal{F}_\mathcal{E}| = k$ where $k \leq n$.

Formally, a serverless choreography configuration for \mathcal{C} is a matrix \mathbf{X} such that:

$$\mathbf{X} \stackrel{def}{=} \{x_{\phi_1}, \dots, x_{\phi_k}\} \in \{\mathbf{F}_{\phi_1} \times \mathbb{N}\} \times \dots \times \{\mathbf{F}_{\phi_k} \times \mathbb{N}\} = \times_{i=1}^k \{\mathbf{F}_{\phi_i} \times \mathbb{N}\} \quad (25)$$

where, for all $i \in [1, k]$, x_{ϕ_i} is an executable function configuration for ϕ_i .

5 The problem of cold starts

In order to find a suitable serverless choreography configuration capable to satisfy QoS imposed by end users, we must find a way to decide

5.1 Metrics

Let $\mathcal{C} = (\Phi, E)$ a basic choreography as defined in section 3.2.4.

Suppose to have a path $\pi = \phi_1 e_1 \phi_2 e_2 \dots e_{n-2} \phi_{n-1} e_{n-1} \phi_n$ and a choreography configuration \mathbf{X} for \mathcal{S} .

We can define the *response time* of π given \mathbf{X} as follows:

$$RT_P(\pi, \mathbf{X}) = \sum_{i=1}^n (RT_{\phi_i}(X(i)) + \sum_{i=1}^{n-1} D(\phi_i, \phi_{i+1})) \quad (26)$$

Similarly, we define the cost of π given \mathbf{X} as follows:

$$C_P(\pi, \mathbf{X}) = \sum_{\substack{1 \leq i \leq n \\ \phi_i \in \mathcal{F}_\varepsilon}} N(\phi_i) \cdot C_{\mathbf{F}_{\phi_i}}(X(i)) \quad (27)$$

Then the response time and the cost of the choreography \mathcal{S} are as follows:

$$RT_C(\mathcal{S}, \mathbf{X}) = \sum_{\pi \in \Pi(\phi_{entry}, \phi_{end})} TPP(\pi) \cdot RT_P(\pi, \mathbf{X}) \quad (28)$$

$$C_C(\mathcal{S}, \mathbf{X}) = \sum_{\pi \in \Pi(\phi_{entry}, \phi_{end})} TPP(\pi) \cdot C_P(\pi, \mathbf{X}) \quad (29)$$

6 Serverless Choreography Structures

Supposing to have a serverless choreography $\mathcal{S} = (\Phi, E)$, we call *structure* any sub-choreography $\mathcal{S}^* = (\Phi^*, E^*)$ whose entry point and the end point are, respectively, an opening and a closing orchestration functions of the same type. Formally:

$$\mathcal{S}^* \text{ is a structure} \Leftrightarrow \begin{cases} \phi_{entry}^* \in \mathcal{F}_{\odot}^* & \wedge & \mathcal{T}(\phi_{entry}^*) = \tau_\alpha \\ \phi_{end}^* \in \mathcal{F}_{\odot}^* & \wedge & \mathcal{T}(\phi_{end}^*) = \tau_\omega \end{cases} \quad (30)$$

The most important aspect is that every structure can be viewed as a set of sub-choreographies of \mathcal{S}^* . Formally, suppose that $\Lambda = (A, B)$ is a graph such that:

$$\begin{aligned} A &\stackrel{def}{=} \Phi^* \setminus \{\phi_{entry}^*, \phi_{end}^*\} \\ B &\stackrel{def}{=} E^* \setminus (out(\phi_{entry}^*) \cup in(\phi_{end}^*)) \end{aligned} \quad (31)$$

Let $c \in \mathbb{N}$ and suppose that c is the number of connected components of the graph Λ .

We denote by $\Theta_{\mathcal{S}^*}$ as the set containing all connected components of Λ every of which is considered as a sub-choreography of \mathcal{S}^* . Formally:

$$\Theta_{\mathcal{S}^*} \stackrel{def}{=} \{\theta_i, \dots, \theta_c\}$$

$$\text{where} \tag{32}$$

$$\begin{aligned} \theta_i &\stackrel{def}{=} (\Phi_i^{**}, E_i^{**}) \text{ is a sub-choreography of } \mathcal{S}^* \\ \Phi_i^{**} &\subset \Phi^* \wedge E_i^{**} \subset E^* \quad \forall i \in [1, c] \end{aligned}$$

6.0.1 Parallel

A structure $\mathcal{P} = (\Phi, E, \Theta)$ is called *parallel structure* when:

$$TPP(\pi) = 1 \quad \forall \pi \in \Pi(\phi_{start}, \phi_{end}) \tag{33}$$

Let $n \in \mathbb{N}$, supposing that $|\Theta| = n$, the response time of a parallel structure is the longest response time of all sub-choreographies $\theta_i \in \Theta$ for all $i \in [1, n]$, while its cost is equal to the sum their costs of execution. Formally, being \mathbf{X} a configuration of \mathcal{P} :

$$RT_{parallel}(\mathcal{P}, \mathbf{X}) = \max \{RT_C(\theta, \mathbf{X}) \mid \theta \in \Theta\} \tag{34}$$

$$C_{parallel}(\mathcal{P}, \mathbf{X}) = \sum_{\theta \in \Theta} C_C(\theta, \mathbf{X}) \tag{35}$$

6.0.2 Branch/Switch

Let $\mathcal{B} = (\Phi, E, \Theta)$ a structure and suppose that $\alpha(\mathcal{B}) = \phi_{entry}$ and $\omega(\mathcal{B}) = \phi_{end}$.

\mathcal{B} is called *branch structure* when following conditions are hold:

$$TPP(\pi) \neq 1 \quad \forall \pi \in \Pi(\phi_{entry}, \phi_{end}) \tag{36}$$

$$|\Theta| = n \quad n \in \{1, 2\} \tag{37}$$

Let all $\theta_i \in \Theta$ sub-choreographies of \mathcal{B} , for all $i \in \{1, 2\}$. Then the response time and the cost of a branch structure can be defined as follows:

$$RT_{branch}(\mathcal{B}, \mathbf{X}) = \sum_{i=1}^n P(\phi_{entry}, \alpha(\theta_i)) RT_C(\theta_i, \mathbf{X}) \tag{38}$$

$$C_{branch}(\mathcal{B}, \mathbf{X}) = \sum_{i=1}^n C_C(\phi_{entry}, \alpha(\theta_i)) C_C(\theta_i, \mathbf{X}) \tag{39}$$

Finally, according to our model, \mathcal{B} is called *switch structure* when $|\Theta| > 2$.

6.0.3 Loop

Let $\mathcal{L} = (\Phi, E, \Theta)$ a structure and suppose that $\alpha(\mathcal{L}) = \phi_{entry}$ and $\omega(\mathcal{L}) = \phi_{end}$.

If $|\Theta| = 1$ and $(\phi_{end}, \phi_{entry}) \in E$, denoting by θ the unique sub-choreography of \mathcal{L} , \mathcal{L} will be called *loop structure* when following conditions are hold:

$$P(\phi_{start}, \alpha(\theta)) = 1 \quad (40)$$

$$P(\phi_{end}, \phi_{entry}) \neq 0 \quad (41)$$

$$P(\phi_{end}, \phi_{entry}) + \sum_{\phi \in succ(\phi_{end})} P(\phi_{end}, \phi) = 1 \quad (42)$$

In order to compute the mean response time and the mean cost of a loop structure, we need to know $EI(\mathcal{L})$, that is the expected value of the number of iterations of \mathcal{L} .

We know that geometric distribution gives the probability that the first occurrence of success requires $k \in \mathbb{N}$ independent trials, each with success probability p and failure probability $q = 1 - p$. In our case, if the our success corresponds to event "we will not execute the loop body", we know that success probability p is given by:

$$p = 1 - P(\phi_{end}, \phi_{entry}) \quad (43)$$

Then:

$$P(EI(\mathcal{L}) = k) = [1 - P(\phi_{end}, \phi_{entry})] \cdot [P(\phi_{end}, \phi_{entry})]^{k-1} \quad (44)$$

$$= pq^{k-1} \quad (45)$$

$$\begin{aligned} E[X] &= \sum_{k=1}^{\infty} (k-1) pq^{k-1} \\ &= p \sum_{k=0}^{\infty} k q^k \\ &= p \cdot \frac{q}{(1-q)^2} \\ &= \frac{q}{p} \\ &= \frac{P(\phi_{end}, \phi_{entry})}{1 - P(\phi_{end}, \phi_{entry})} \end{aligned} \quad (46)$$

Then the response time and the cost of a loop structure can be defined as follows:

$$RT_{loop}(\mathcal{L}, \mathbf{X}) = E[I_{\mathcal{L}}] \cdot RT_C(\theta, \mathbf{X}) \quad (47)$$

$$C_{loop}(\mathcal{L}, \mathbf{X}) = E[I_{\mathcal{L}}] \cdot C_C(\theta, \mathbf{X}) \quad (48)$$

7 Optimization

7.1 Performance Modeling

We define the utility function as follows:

$$F(\mathbf{x}) \stackrel{def}{=} w_{RT} \cdot \frac{RT_{max} - RT(\mathbf{x})}{RT_{max} - RT_{min}} + w_C \cdot \frac{C_{max} - C(\mathbf{x})}{C_{max} - C_{min}} \quad (49)$$

where $w_{RT}, w_C \geq 0$, $w_{RT} + w_C = 1$, are weights for the different QoS attributes, while RT_{max} (RT_{min}) and C_{max} (C_{min}) denote, respectively, the maximum (minimum) value for the overall expected response time and cost.

$$\arg \max_{\mathbf{X}} F(\mathbf{X})$$

$$\text{subject to } C(\mathbf{X}) \leq C_{user}$$

$$RT(\mathbf{X}) \leq RT_{user} \quad (50)$$

$$\sum_{j=1}^{|\mathbf{F}_{\phi_i} \times \mathbb{N}|} x_{\phi_{ij}} = 1 \quad \forall \phi_i \in \mathcal{F}_{\mathcal{E}}$$

$$x_{\phi_{ij}} \in \{0, 1\} \quad \forall \phi_i \in \mathcal{F}_{\mathcal{E}} \wedge j \in \mathbb{N} \cap [1, |\mathbf{F}_{\phi_i} \times \mathbb{N}|]$$

fdsfsdfsdfs

$$\sum_{i=0}^{|\mathcal{F}_{\mathcal{E}}|} \sum_{j=0}^{|\mathbf{F}_{\phi_i}|} \sum_{s \in M \subset \mathbb{N}} F(x_{\phi_0}, \dots, x_{\phi_{|\mathcal{F}_{\mathcal{E}}|}}) \quad (51)$$

8 Super-peer Node

An *end-user* represents, instead, a third-party application that wants execute one or more function choreographies.

9 Peer Node

Proposed

Hybrid structures are notably deployed in collaborative distributed systems. The main issue in many of these systems is to first get started, for which often a traditional client-server scheme is deployed. Once a node has joined the system, it can use a fully decentralized scheme for collaboration.

Relating to a specified function choreography X belonging to resource owner R , a peer P of our system can be in one of the following states:

Active State When P has been marked as responsible for manage all invocation requests of X forwarded by end users.

Forwarder State Otherwise

function choreographies (FCs) or workflows of functions.

As known, in server-less computing platforms, computation is done in **function instances**. These instances are completely managed by the server-less computing platform provider (SSP) and act as tiny servers where a function is been executed.

10 Resources

11 System's resources and actors

Given a resource owner R , there are two type of resources which he can manage:

1. Function choreographies.
2. Server-less function implementations (also called *concrete server-less functions*)

11.1 Server-less function swarms

Informally, a so-called *server-less function swarm* represent a set of concrete server-less function with very specific properties.

Precisely, let $l \in \mathbb{N}$ such that $l \neq 0$, R a resource owner, P a server-less computing platform provider and X_R a set of concrete server-less functions. Moreover, let \mathbf{X}_R the set containing all concrete function implementations defined and deployed by R on any provider.

A set $X_R \subseteq \mathbf{X}_R$ is called a *server-less function swarm*, or simply *swarm*, if:

1. $|X_R| = n \geq 1$, that is X_R must contain at least one concrete function.
2. X_R contain concrete functions that share the same platform provider P where they will be executed.
3. X_R contain concrete functions that share the same limit l in term of max number of server-less function instance runnable at the same time by the corresponding platform provider P . That limit is also called *server-less function swarm's concurrency limit*, or simply, *concurrency limit*.

Is very important to make clear that only at most l concrete functions belonging to X_R can be executed simultaneously by P . The value of l depends by specific policies adopted by P ; some of them imposed that limit *per-account*, others *per-functions*. Our model supports both approaches because:

- If P imposed limits *per-function*, then $|X_R| = 1$, that is, X_R will contain only one function defined and deployed by R in P , where l will be represent the provider's per-function limit.
- If P imposed a limit *per-account*, then generally $|X_R| \geq 1$ and it include all concrete server-less function deployed on P while l will be represent the provider's global limit.

11.1.1 Server-less function sub-swarms

A sub-swarm of X_R , which we will denote with Δ_{X_R} , is the term with which we denote any element belong to the power set¹ of X_R , excluding the empty set. Formally, any $\Delta_{X_R} \in \mathcal{P}(X_R) \setminus \emptyset$ is a sub-swarm.

Is very important to remember that in our model any sub-swarm Δ_{X_R} of X_R has the *same* concurrency limit of X_R .

¹The *power set* $\mathcal{P}(S)$ of a set S is the set of all subsets of S , including the empty set and S itself.

12 Function choreography scheduler

12.1 Schedulability condition

Let FC_R a function choreography belonging to a resource owner R and $F_{abstract}$ its server-less abstract functions set. Moreover, be \mathbf{X}_R the set containing all functions deployed by R in any provider.

In order to effectively start the execution of a function choreography, is required that for each abstract function $f_{abstract} \in F$ at least one concrete function f , which implements it, exists.

Formally, a function choreography is said *schedulable* when:

$$FC_R \text{ is schedulable} \Leftrightarrow \begin{aligned} & \forall f_{abstract} \in F_{abstract} \\ & \exists f \in \mathbf{X}_R \mid f \text{ implements } f_{abstract} \end{aligned} \quad (52)$$

Although it is correct, the condition expressed by equation 52 is not very precise, because \mathbf{X}_R can contain some functions that doesn't implement any $f_{abstract} \in F_{abstract}$.

Therefore, we define Ω_{FC_R} as the set containing only concrete functions that are needed to execute FC_R , which can both to belong to any provider and to have different concurrency limits.

Since multiple implementations of a same abstract function can exist at the same time, we can exploit the notion of swarm and sub-swarm to formally define the set Ω_{FC_R} .

Let $n \in \mathbb{N}$, such that $n \geq 1$, and X_{R_i} the i -th swarm and $\Delta_{X_{R_i}}$ its sub-swarm which contains only concrete functions implementing one or more $f_{abstract} \in F_{abstract}$, where $1 \leq i \leq n$.

We define Ω_{FC_R} as follows:

$$\begin{aligned} \Omega_{FC_R} &\stackrel{def}{=} \Delta_{X_{R_1}} \cup \dots \cup \Delta_{X_{R_n}} = \bigcup_{i=1}^n \Delta_{X_{R_i}} \\ &\text{where} \\ \Delta_{X_{R_i}} \cap \Delta_{X_{R_j}} &= \emptyset \quad \text{for } i, j \in [0, n] \mid i \neq j \\ \forall f \in \Delta_{X_{R_s}} & \quad f \text{ implements } f_{abstract} \text{ for } s \in [0, n] \end{aligned} \quad (53)$$

Since belong to different swarms, please note that any $\Delta_{X_{R_i}}$ and $\Delta_{X_{R_j}}$, for any $i \neq j$, can belong to the same provider but they cannot share the same concurrency limit.

Generally the schedulability condition for FC_R can be written as follows:

$$FC_R \text{ is schedulable} \Leftrightarrow \exists \Omega_{FC_R} \quad (54)$$

12.2 The Δ_{X_R} -Scheduler

Let R a resource owner, P the server-less provider and Δ_{X_R} a sub-swarm of a X_R , where k its concurrency limit.

Let $m \in \mathbb{N}$, a Δ_{X_R} -Scheduler, denoted as $S_{(\Delta_{X_R}, m)}$ represents a queuing system, implementing any scheduling discipline, equipped with m so-called *virtual function instance*, where $m \leq k$.

Its aim is to decide when and which function, belonging to Δ_{X_R} , must be performed on P .

The parameter m is also called *scheduler capacity*.

12.2.1 Virtual function instance

A *virtual function instance* represents a real function instances, clearly belonging to the server-less computing platform provider, which is *virtually* owned by $S_{(\Delta_{X_R}, m)}$.

Therefore, m represents the max number of server-less function instances usable simultaneously by $S_{(\Delta_{X_R}, m)}$.

12.2.2 Proprieties and constrains

According to our model, a Δ_{X_R} -scheduler capable to manage any function belonging to Δ_{X_R} , if exist, is *not* unique, although it is unique inside a peer node.

In order to achieve better performance in terms of network delay experienced by end users, fault tolerance and load balance, any peer nodes can hold a Δ_{X_R} -scheduler in order to manage incoming request sent by several users spread in different geographic regions.

However, despite there is no upper bound to the number of Δ_{X_R} -schedulers existing at the same time in our system, there is a limitation regarding the scheduler capacity of each existing scheduler.

Let's start summarizing all rules regarding Δ_{X_R} -schedulers:

1. All peer node of our system can hold a Δ_{X_R} -scheduler object.
2. Each node can hold only one instance of type Δ_{X_R} -scheduler.
3. Let $n \in \mathbb{N}$ such that $n \geq 1$, suppose that our system contains n peer nodes holding a Δ_{X_R} -scheduler.

To be more precise, let's say that a sequence $S_{(1, (\Delta_{X_R}, m_1))}, \dots, S_{(i, (\Delta_{X_R}, m_i))}$ exist at the same time in our system, where $S_{(i, (\Delta_{X_R}, m_i))}$ represent the Δ_{X_R} -scheduler owned by i -th node having scheduler capacity equal to m_i .

Following constraint must be hold:

$$\sum_{i=1}^n m_i \leq k \quad (55)$$

where $k \in \mathbb{N}$, with $k > 0$, is the concurrency limit of the swarm X_R .

Remember that any sub-swarm Δ_{X_R} share the same concurrency limit of X_R . Therefore, equation 59 states that, the sum of all scheduler capacities which manage the functions belonging to Δ_{X_R} , must be less or equal to the max number of function instances executable at the same time on the server-less computing platform provider.

12.3 The FC_R -Scheduler

To support hybrid-scheduling, that is the ability to execute multiple concrete function implementations belonging to different providers or subjected to different concurrency limit, in order to select the most suitable concrete function

implementation according to a given QoS, unfortunately only one “*scheduler*” is not enough.

We call *FC_R-Scheduler* a set of *Δ_{X_R}-Schedulers* where $\Delta_{X_R} \in \Omega_{FC_R}$. Is always required that $|FC_R| \geq 1$, that is, at least one scheduler must exist.

13 FC_R-Active Peer Node

According to our model, in order to effectively invoke all server-less concrete function belonging to a function choreography *FC_R*, is required that a peer node is “active”.

We said that a peer node *A* is *FC_R-active peer node*, or, simply, *active*, when it holds a *Δ_{X_R}-scheduler* for any sub-swarm in Ω_{FC_R} . Formally:

$$A \text{ is } FC_R\text{-active peer node} \Leftrightarrow \forall \Delta_{X_R} \in \Omega_{FC_R} \quad \exists S_{(\Delta_{X_R}, m)} \text{ hold by } A \quad (56)$$

Multiple nodes can be active at the same time. Any node perform its scheduling decision independently.

14 Architecture overview

Our system design is based on a network of nodes, or *peers*, every of which has the *same functionality*; in fact, any of them is able to handle request submission, request scheduling and, potentially, request execution.

This is the reason according to which we can mark our proposal as a *P2P system*.

14.1 Overlay network

Our system’s nodes are connected by an *overlay network*.

Definition 14.1 (Overlay network). An overlay network is a virtual network built on top of a physical network according to which each nodes can communicate with other if and only if they are connected by virtual links belonging to the virtual network.

Remark. A node may not be able to communicate directly with an arbitrary other node although they can communicate through physical network.

To be more precise, we have adopt a fully *centralized unstructured overlay network* because the *peer-resource index*, sometimes called *directory*, is centralized.

Please note that hybrid unstructured or fully decentralized solutions are technically possible, but guarantees about quality of service are very difficult.

14.1.1 Locality-awareness property

Locality-awareness is one of the essential characteristics of our system. In fact, if each peer is able to select his neighbours exploiting a suitable locality aware algorithm, is possible to decrease user experienced delays.

We have decided to adopt a multi-level based locality-aware neighbour selection called *intra-AS lowest delay clustering algorithm* (ASLDC).

When ASLDC algorithm is used, each peer chooses nearby peers only from those within the same AS; then it ranks its neighbours in terms of transmission latency, preferring to establish the connection with the node with the shortest latency to itself.

TODO [3]

14.2 Fault Tolerance

Like in many other P2P implementation, there are two ways of detecting failures in our proposed solution:

1. If a node tries to communicate with a neighbour and fails.
2. Since all nodes send to all his neighbour nodes so-called “heartbeat” messages, that is messages sent at fixed time intervals to indicate that the sender is alive, is possible to detect a failure by not receiving aforementioned periodic update messages after a long time.

14.2.1 Availability

Let $n \in \mathbb{N}$, R a resource owner and FC_R his function choreography, since multiple FC_R -Scheduler object can exist in n different nodes, our proposed solution can guarantee an high degree of availability.

In fact, when a node holding an FC_R -Scheduler fails, all end users requests related to FC_R can be routed to any other node holding a FC_R -Scheduler.

14.2.2 Replication

Although multiple FC_R -Scheduler object can exist, it's not mean that FC_R -Scheduler object are replicated, because every of them manage different virtual function instances.

Except the peer-resource index, any form of replication is performed by our system.

Algorithm 1 An algorithm with caption

$y \leftarrow 1$

and Δ_{X_R} -Scheduler

The locality awareness of the overlay network is used in scheduling jobs, described in Section 3.2

physical network but with added properties such as fault tolerance and flexibility.

The first problem is scheduling. Since there is no central scheduler it is difficult

Hybrid unstructured overlay

14.3 FC_R -Request

object with only one limitation:

Suppose that globally there are a set of schedulers $S_{1,(R_X,m_1)}, \dots, S_{p,(R_X,m_p)}$, where $p \in \mathbb{N}$ with $p \geq 1$

To be more precise, when a function x_j must to be execute, let s the current number of busy virtual instances, one of the following events may occur:

1. if $s < m$, the scheduler invoke directly the function x_j on the provider.
2. if $s = m$, the scheduler delay the execution of the function x_j on the provider according to implemented scheduling discipline.

Let R a resource owner and R_x its function choreography made up of $R_{x_1}, R_{x_2}, \dots, R_{x_n}$ unique server-less functions; it is said that a peer node P is **responsible** for R_x when it contains a sequence of schedulers $S_{R_1}, S_{R_2}, \dots, S_{R_k}$ with $k \leq n$, belonging to R , capable to invoke all server-less function belonging to R_x .

It is said that a

Depending on the definition of the function choreography provided by R and the unique characteristics of back-end server-less providers which execute all serverless functions R_{x_n} of

It is said that a scheduler S is capable to invoke a server-less function when , a scheduler S can invoke multiple

When a peer A , placed “*at the edge*” of the network, receives a new request of invocation for X by an end user, it performs following task in that order:

1. If it responsible It check for it is an already an *active peer* to manage

has found the tracker for a file F , the tracker returns a subset of all the nodes currently involved in downloading F .

Replication and Fault Tolerance. There are two ways of detecting failures in CAN, the first if a node tries to communicate with a neighbor and fails, it takes over that neighbor’s zone. The second way of detecting a failure is by not receiving the periodic update message after a long time. In the second case, the failure would probably be detected by all the neighbors, and all of them would try to take over the zone of the failed node, to resolve this, all nodes send to all other neighbors the size of their zone, and the node with the smallest zone takes over.

$$E[T] = \sum_{i=0}^n E[S_i] + E[T_{Q_i}] \quad (57)$$

15 Queuing system of a Δ_{X_R} -scheduler

Let $k \in \mathbb{N}$ such that $k \geq 0$. Moreover, suppose that R represents a resource owner, P a server-less provider and Δ_{X_R} a sub-swarm of the swarm X_R where k its concurrency limit. Finally, let $S_{(\Delta_{X_R}, m)}$ a Δ_{X_R} -scheduler.

Obliviously, since there are only m available virtual function instances, if there are more than m server-less functions waiting to be execute, a choice has to be made about which server-less function has to run next to ensuring QoS guarantees for latency critical applications.

Because we expect to execute server-less functions with different response-time requirements, which may have different scheduling needs, i have designed the Δ_{X_R} -scheduler as a queuing system implementing a *multilevel queue scheduling algorithm* which partitions the ready queue into several separate queues.

According to our solution, each queue has its own scheduling algorithm and any server-less function is permanently assigned to one queue according to his class.

In addition, we have adopt *round-robin* (RR) scheduling algorithm to perform scheduling activity among the queues.

15.1 Server-less function preemption

A very important consideration regards server-less function preemption.

In our context, due to FaaS paradigm, which hides the complexity of servers where our functions will be executed, the ability to preempt functions is not naturally available.

For that reason, only non-preemptive scheduling algorithms, according to which once a server-less function starts running it cannot be preempted, even if a higher priority server-less function comes along, can be adopted.

Since many scheduling algorithms require job preemption to run optimally, this situation can lead to a suboptimal resource management.

15.2 Queuing system design

A Δ_{X_R} -scheduler is made up of three queues:

- One queue implements a *Non-Preemptive Least-Slack-Time-First* (LST) scheduling algorithm. We will refer to that queue using Q_{LST} notation.

That algorithm implements a *dynamic priority scheduling approach* where priorities are assigned to server-less functions based on their *slacks*.

At any time t , the *slack* of a job with deadline at d is equal to $d - t - s$, where s is the time required to complete the job.

Any job having strict latency requirements must to be assign to this queue.

- Another queue implements, instead, a *Non-Preemptive Shortest-Job-First* (*SJF*) scheduling algorithm. That queue is denoted with Q_{SJF} .

That policy assigns priorities to jobs based on their size: the smaller the size, the higher the priority.

This queue is reserved for any function choreography that tolerates high latency. However, according to queueing theory, since SJF performs very bad when service time distribution is heavy-tailed; in fact

$$E[T_Q(x)]^{SJF} = \frac{\rho E[S^2]}{2E[S]} \cdot \frac{1}{(1 - \rho x)^2} \quad (58)$$

This queue is not suitable for very, very large server-less function. The variance in the job size distribution must be low.

- Finally, the third queue exploits the the simplest scheduling algorithm, that is the *First-Come-First-Served* (*FCFS*) policy, and it is denoted using Q_{FCFS} notation.

This queue is used for any heavy-tailed server-less function which can tolerate high latency.

15.3 The virtual function instances allocation problem

Let $n \in \mathbb{N}$ such that $n \geq 1$ and suppose that, at a certain time, a sequence including n unique Δ_{X_R} -schedulers exist in our system.

Formally, that sequence is denoted with $S_{(1,(\Delta_{X_R},m_1))}, \dots, S_{(i,(\Delta_{X_R},m_n))}$, where $S_{(i,(\Delta_{X_R},m_i))}$ represents the Δ_{X_R} -scheduler owned by i -th node having scheduler capacity equal to m_i .

We have already established that the following constraint must be hold:

$$\sum_{i=1}^n m_i \leq k \quad (59)$$

where $k \in \mathbb{N}$, with $k > 0$, is the concurrency limit of the swarm X_R .

15.3.1 Reactive scaling policy

Every node of the system monitors all queues belong to its Δ_{X_R} -Scheduler and adjusts the number of reserved virtual function instances to preserve quality-of-service guarantees. This activity is called *reactive scaling policy*.

Although a $S_{(\Delta_{X_R},m)}$ owns m virtual function instances, they are not equally shared between aforementioned run queue.

First of all, let be v_{LST} , v_{SJF} and v_{FCFS} the numbers of the virtual function instances bounded to the Q_{LST} , Q_{SJF} and Q_{FCFS} run-queue.

At any time, following constrains must be hold:

$$v_{LST} \stackrel{def}{=} \frac{N \cdot E[S]}{\sum slack} \quad (60)$$

$$v_{SJF} \stackrel{def}{=} \frac{len(Q_{SJF})}{\alpha} \quad (61)$$

$$v_{FCFS} \stackrel{def}{=} \frac{len(Q_{FCFS})}{\alpha} \quad (62)$$

$$m = v_{LST} + v_{SJF} + v_{FCFS} \quad (63)$$

1. If the current number of virtual function instances is not enough to avoid deadline misses by latency constrained server-less function inside Q_{LST} , will be add to $VCPU_{LST}$ a number of virtual processor equal to

$$v_{LST} < \frac{N \cdot E[S]}{\sum slack} \quad (64)$$

This form of aging prevents starvation.

Finally, suppose that a node peer A holds a $S_{(\Delta_{X_R}, m_A)}$, which represents the Δ_{X_R} -Scheduler object.

1. Once a request is received over HTTP
2. Ask to the coordinator for

$$m = \frac{\sum slack}{E[S] \cdot m_{current}} + \frac{N}{\alpha} \quad (65)$$

$$\lambda_{FC_R} \geq \lambda_{threshold} \quad (66)$$

$$RTT_{(A,B)} > RTT_{(A,C)} + \quad (67)$$

$$f_{abstract} = next \quad (68)$$

$$f_{abstract} = min(E[f_{\Delta_{X_{R_i}}}] + E[T_{Q_{LST}}]F \quad (69)$$

15.3.2 Server-less function scheduling for QoS “Minimum Response Time”

Let $n \in \mathbb{N}$, such that $n \geq 1$, supposing to have an FC_R -active peer node and let Ω_{FC_R} the set of all sub-swarms containing at least one function implementing any $f_{abstract} \in \mathbf{F}_{abstract}$, where $|\Omega_{FC_R}| = n$.

Finally, suppose that $f_{abstract}$ is the next abstract function that have to be executed according to the control-flow logic described by FC_R .

1. For $1 \leq i \leq n$, let $\Delta_{(X_{R_i}, f_{abstract})} \subseteq \Delta_{X_{R_i}}$ containing all concrete functions implementing $f_{abstract}$.
2. For $1 \leq i \leq n$, select $\mathbf{f}_i \in \Delta_{(X_{R_i}, f_{abstract})}$ functions having the minimum response time $E[T_{\mathbf{f}_i}]$.

3. Found $i \in [0, m]$ such that:

$$E[T_{\mathbf{f}_i}] + E[T_{Q_i}]^{\mathbf{LST}} < E[T_{\mathbf{f}_j}] + E[T_{Q_j}]^{\mathbf{LST}} \quad \text{for } j \in \mathbb{N}, 1 \leq j \leq n, i \neq j \quad (70)$$

4. Select the i -th Δ_{X_R} -Schedulers belonging to FC_R -Scheduler of the node and add \mathbf{f}_i into the Q_{LST} run-queue.