

References

- [1] Nabeel Akhtar, Ali Raza, Vatche Ishakian, and Ibrahim Matta. Cose: Configuring serverless functions using statistical learning. In *IEEE INFOCOM 2020 - IEEE Conference on Computer Communications*, pages 129–138, July 2020.
- [2] Nima Mahmoudi and Hamzeh Khazaei. Performance modeling of serverless computing platforms. *IEEE Transactions on Cloud Computing*, pages 1–1, 2020.
- [3] Xiongfei Weng, Hongliang Yu, Guangyu Shi, Jian Chen, Xu Wang, Jing Sun, and Weimin Zheng. Understanding locality-awareness in peer-to-peer systems. In *2008 International Conference on Parallel Processing - Workshops*, pages 59–66, 2008.

Contents

1	Serverless Computing Paradigm	4
1.1	Serverless Function	4
1.2	Serverless Application	4
2	Orchestrator Model	4
2.0.1	Abstract server-less function	5
2.0.2	Concrete server-less function	5
2.0.3	Control-flow operator	6
3	Performance Modeling	6
3.1	Definition of serverless workflow	6
4	Aim	7
5	Super-peer Node	8
6	Peer Node	8
7	Resources	8
8	System’s resources and actors	8
8.1	Resource owner	8
8.2	Server-less function swarms	8
8.2.1	Server-less function sub-swarms	9
9	Function choreography scheduler	9
9.1	Schedulability condition	9
9.2	The Δ_{X_R} -Scheduler	10
9.2.1	Virtual function instance	10
9.2.2	Proprieties and constrains	10
9.3	The FC_R -Scheduler	11
10	FC_R-Active Peer Node	11

11 Architecture overview	12
11.1 Overlay network	12
11.1.1 Locality-awareness property	12
11.2 Fault Tolerance	12
11.2.1 Availability	13
11.2.2 Replication	13
11.3 FC_R -Request	13
12 Queuing system of a Δ_{X_R}-scheduler	15
12.1 Server-less function preemption	15
12.2 Queuing system design	15
12.3 The virtual function instances allocation problem	16
12.3.1 Reactive scaling policy	16
12.3.2 Server-less function scheduling for QoS “Minimum Response Time”	17

fdsfsfsdfsdfs f sd fsd f sdfsd fsdfsdsdterterterterterterterterterterter

$$\begin{aligned} \min \quad & \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} g(x)^{f_k^i} Y_x^{f_k^i} \\ & \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} T(x)^{f_k^i} Y_x^{f_k^i} \leq D_{\mathbf{F}} \\ & \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} Y_x^{f_k^i} = 1 \end{aligned} \quad (1)$$

$$Y_x^{f_k^i} \in \{0, 1\}$$

$$\begin{aligned} \min \quad & \sum_{\pi \in \Pi} \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} g(x)^{f_k^i} Y_{(\pi, x)}^{f_k^i} \\ & \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} T(x)^{f_k^i} Y_x^{f_k^i} \leq D_{\mathbf{F}} \\ & \sum_{\pi \in \Pi} \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} Y_{(\pi, x)}^{f_k^i} = 1 \end{aligned} \quad (2)$$

$$Y_{(\pi, x)}^{f_k^i} \in \{0, 1\}$$

1 Serverless Computing Paradigm

In serverless computing platforms, computation is done by so-called *function instances* which are completely managed by the serverless computing platform provider and act as tiny servers that can be invoked based on events forwarded by end-users [2].

Serverless computing platforms handle almost every aspect of the system administration tasks needed to deploy a workload on the cloud, providing a new simplified programming model according to which developers can focus on the business aspects of their applications only [1].

Moreover, the paradigm lowers the cost of deploying applications too, adopting a so-called “*pay as you go*” pricing model, by charging for execution time rather than for allocated resources [1].

1.1 Serverless Function

A *serverless function* represents a stateless, event-driven, self-contained unit of computation implementing a business functionality.

Although a serverless function generally represents a unit of executable code, submitted to FaaS platforms by developers using one or a combination of the programming languages supported by FaaS providers, a serverless function can be any cloud services eventually necessary to business logic, like cloud storage, message queue services, pub/sub messaging service etc.

When it represents executable code, developers can specify several configuration parameters, like timeout, memory size or CPU power [1].

A serverless function can be triggered through events or HTTP requests following which the FaaS provider executes it in a containerized environment, like container, virtual machine or even processes, using the specified configuration.

1.2 Serverless Application

A *serverless application* represents a stateless and event-driven software system made up of a set of serverless functions hosted on one or more FaaS platforms and combined together by an *orchestrator*, which handles events among functions and triggers them in the correct order according to the business logic defined by developers.

2 Orchestrator Model

2.1 Resource owner

A *resource owner*, henceforward denoted with R , represents an entity capable of *creating*, *modifying* and *authorizing* access to several resources of our system.

A *function choreography* is the most important resource of our system and it is used to model both server-less functions and server-less function compositions.

Informally, it represents, using graph notation, all paths that might be traversed through a server-less function composition during its execution. In other words, it describes calling relationships between server-less functions.

Formally, being R the resource owner, a **function choreography**, denoted as FC_R , is a *control-flow graph* $G(V, E)$, where:

- Each node $v \in V$, called *abstract control-flow node*, represents a generic function of a computer program.
- Each edge $(v_i, v_j) \in E$, for any $i, j \in \mathbb{N}$ with $i \neq j$, indicates that the generic function v_i calls that in v_j .

Any function choreography can be uniquely identified by an ordered string pair (a, b) , where a is the *resource owner name* while b is the *function choreography name*.

There are two types of abstract control-flow nodes:

Abstract server-less function denoted as $f_{(abstract,i)}$ while the set of all server-less functions is denoted as $F_{abstract}$.

Control-flow function denoted as c_t while C is the corresponding set of all control-flow functions.

Formally, let $|V| = n$, $|F_{abstract}| = t$, and $|C| = k$, where $t, k \in \mathbb{N}$ with $t + k = n$, we have that:

$$\begin{aligned} V &\stackrel{def}{=} F_{abstract} \cup C \\ &= (f_{(abstract,1)}, \dots, f_{(abstract,t)}) \cup (c_1, \dots, c_k) \end{aligned} \quad (3)$$

Clearly, we said that function choreography models a server-less function, when $|V| = 1$ and $|E| = 0$; conversely, it models a server-less function composition, when $|V| > 1$ and $|E| > 0$.

2.1.1 Abstract server-less function

An *abstract server-less function* represents a descriptions of one or more corresponding concrete server-less function implementations. That description includes:

- TODO
- TODO
- TODO

Any abstract server-less function can be uniquely identified by an ordered string pair (a, b) , where a is the *resource owner name* while b is the *abstract server-less function name*.

2.1.2 Concrete server-less function

Given an abstract server-less function, a resource owner can provide different implementations which, although they must be semantically equivalent, may eventually expose different performance or cost behaviour.

Therefore, we call *concrete server-less function* any implementation of a given abstract function and it is uniquely identified by an ordered string tuple (a, b, c) , where a and b represent, like before, the *resource owner name* and the *abstract function name* respectively, while c represents the so-called *function type*, which is an abstract descriptions of the corresponding function implementation.

2.1.3 Control-flow operator

Like `if`, `for`, etc.

3 Performance Modeling

3.1 Definition of serverless workflow

Let $n \in \mathbb{N} \setminus \{0\}$, a serverless workflow G_s is defined as a weighted directed graph:

$$G_s \stackrel{def}{=} (V, E, P, D, RT, RTTP, M, NI, C) \quad (4)$$

where

- V represents a finite set of $|V| = n$ vertices such that $V = \{\mathcal{F}_1, \mathcal{F}_2, \dots, \mathcal{F}_{n-1}, \mathcal{F}_n\}$, where, for all $i \in [1, n]$, each \mathcal{F}_i represents an abstract FaaS function;
- $E \subseteq V \times V$ represents a finite set of directed edges where, for all $v, u \in [1, n]$, each directed edge from $\mathcal{F}_u \in V$ to $\mathcal{F}_v \in V$, denoted as $(\mathcal{F}_u, \mathcal{F}_v)$, represents the interaction between \mathcal{F}_u and \mathcal{F}_v defined by the business logic;
- $P : V \times V \rightarrow [0, 1]$ is a the transition probability function where $P(\mathcal{F}_u, \mathcal{F}_v)$ represents the probability of invoking \mathcal{F}_v after finishing the execution of \mathcal{F}_u .

Clearly, when $P(\mathcal{F}_u, \mathcal{F}_v) = 0$, it means that the $(\mathcal{F}_u, \mathcal{F}_v)$ does not exist.

4 Aim

Let G_s a serverless workflow, the basic goal of our framework is to determine a so-called *serverless workflow configuration* π that allows us to meet user-specified QoS objectives.

Precisely, \mathbf{x} is a vector such that

$$\mathbf{x} \in \{\mathbf{F}_1 \times \mathbb{N}\} \times \dots \times \{\mathbf{F}_n \times \mathbb{N}\} = \bigtimes_{i=1}^n \{\mathbf{F}_i \times \mathbb{N}\} \quad (5)$$

In other words, for all $1 \leq i \leq n$, \mathbf{x} is a vector containing n ordered pairs (f_i, m_i) called *configuration* of the abstract serverless function \mathcal{F}_i , where f_i represents a concrete serverless function while m_i denotes its allocated memory size. In other words, the aim of our framework is to find the most suitable configuration in order to achieve user-specified QoS objectives.

A path in the serverless workflow is a finite sequence of distinct vertices and edges $w = \mathcal{F}_1 e_1 \mathcal{F}_2 e_2 \dots \mathcal{F}_{n-1} \mathcal{F}_n$ such that:

The transition probability of the simple path w

Before to proceed, when the serverless workflow configuration x is used,

- $D(\mathbf{x})$ denotes the end-to-end response time of current serverless workflow G_s when the serverless workflow configuration x is used. At the same time, $C(\mathbf{x})$ represents the

- $C(\mathbf{x})$ re

gdfgdfgdfgd

- $\mathbf{x}(\mathcal{F})$ is used to denote the configuration (f, m) of the serverless abstract function \mathcal{F} .
- $C_{func}(\mathbf{x}(\mathcal{F}))$ represents the cost to pay to execute the concrete serverless function f with memory size m .

$$TPP(w) = \prod_{i=1}^{n-1} P(\mathcal{F}_i, \mathcal{F}_{i+1}) \quad (6)$$

$$C_{walk}(w, \mathbf{x}) = \sum_{\mathcal{F} \in w} C_{func}(\mathbf{x}(\mathcal{F})) \quad (7)$$

$$C_{conf}(\mathbf{x}) = \sum_{w \in SP_{G_s}(\mathcal{F}_{start}, \mathcal{F}_{end})} TPP(w) C(w, \mathbf{x}) \quad (8)$$

$$\arg \min_{\pi} \sum_{i=1}^n C_{\pi}(x_{ij}) x_{ij}$$

$$\begin{aligned} \sum_{n=1}^N \sum_{k=1}^K \sum_{x \in C} T(x)^{f_k^i} Y_x^{f_k^i} &\leq D_{\mathbf{F}} \\ \sum_{j=1}^{n_i} x_{ij} &= 1 \quad \forall i \in \{0, n\} \end{aligned} \quad (9)$$

$$x_{ij} \in \{0, 1\}$$

5 Super-peer Node

An *end-user* represents, instead, a third-party application that wants execute one or more function choreographies.

6 Peer Node

Proposed

Hybrid structures are notably deployed in collaborative distributed systems. The main issue in many of these systems is to first get started, for which often a traditional client-server scheme is deployed. Once a node has joined the system, it can use a fully decentralized scheme for collaboration.

Relating to a specified function choreography X belonging to resource owner R , a peer P of our system can be in one of the following states:

Active State When P has been marked as responsible for manage all invocation requests of X forwarded by end users.

Forwarder State Otherwise

function choreographies (FCs) or workflows of functions.

As known, in server-less computing platforms, computation is done in **function instances**. These instances are completely managed by the server-less computing platform provider (SSP) and act as tiny servers where a function is been executed.

7 Resources

8 System's resources and actors

Given a resource owner R , there are two type of resources which he can manage:

1. Function choreographies.
2. Server-less function implementations (also called *concrete server-less functions*)

8.1 Server-less function swarms

Informally, a so-called *server-less function swarm* represent a set of concrete server-less function with very specific properties.

Precisely, let $l \in \mathbb{N}$ such that $l \neq 0$, R a resource owner, P a server-less computing platform provider and X_R a set of concrete server-less functions. Moreover, let \mathbf{X}_R the set containing all concrete function implementations defined and deployed by R on any provider.

A set $X_R \subseteq \mathbf{X}_R$ is called a *server-less function swarm*, or simply *swarm*, if:

1. $|X_R| = n \geq 1$, that is X_R must contain at least one concrete function.
2. X_R contain concrete functions that share the same platform provider P where they will be executed.

3. X_R contain concrete functions that share the same limit l in term of max number of server-less function instance runnable at the same time by the corresponding platform provider P . That limit is also called *server-less function swarm's concurrency limit*, or simply, *concurrency limit*.

Is very important to make clear that only at most l concrete functions belonging to X_R can be executed simultaneously by P . The value of l depends by specific policies adopted by P ; some of them imposed that limit *per-account*, others *per-functions*. Our model supports both approaches because:

- If P imposed limits *per-function*, then $|X_R| = 1$, that is, X_R will contain only one function defined and deployed by R in P , where l will be represent the provider's per-function limit.
- If P imposed a limit *per-account*, then generally $|X_R| \geq 1$ and it include all concrete server-less function deployed on P while l will be represent the provider's global limit.

8.1.1 Server-less function sub-swarms

A sub-swarm of X_R , which we will denote with Δ_{X_R} , is the term with which we denote any element belong to the power set¹ of X_R , excluding the empty set. Formally, any $\Delta_{X_R} \in \mathcal{P}(X_R) \setminus \emptyset$ is a sub-swarm.

Is very important to remember that in our model any sub-swarm Δ_{X_R} of X_R has the *same* concurrency limit of X_R .

9 Function choreography scheduler

9.1 Schedulability condition

Let FC_R a function choreography belonging to a resource owner R and $F_{abstract}$ its server-less abstract functions set. Moreover, be \mathbf{X}_R the set containing all functions deployed by R in any provider.

In order to effectively start the execution of a function choreography, is required that for each abstract function $f_{abstract} \in F$ at least one concrete function f , which implements it, exists.

Formally, a function choreography is said *schedulable* when:

$$FC_R \text{ is schedulable} \Leftrightarrow \begin{aligned} &\forall f_{abstract} \in F_{abstract} \\ &\exists f \in \mathbf{X}_R \mid f \text{ implements } f_{abstract} \end{aligned} \quad (10)$$

Although it is correct, the condition expressed by equation 10 is not very precise, because \mathbf{X}_R can contain some functions that doesn't implement any $f_{abstract} \in F_{abstract}$.

Therefore, we define Ω_{FC_R} as the set containing only concrete functions that are needed to execute FC_R , which can both to belong to any provider and to have different concurrency limits.

¹The *power set* $\mathcal{P}(S)$ of a set S is the set of all subsets of S , including the empty set and S itself.

Since multiple implementations of a same abstract function can exist at the same time, we can exploit the notion of swarm and sub-swarm to formally define the set Ω_{FC_R} .

Let $n \in \mathbb{N}$, such that $n \geq 1$, and X_{R_i} the i -th swarm and $\Delta_{X_{R_i}}$ its sub-swarm which contains only concrete functions implementing one or more $f_{abstract} \in F_{abstract}$, where $1 \leq i \leq n$.

We define Ω_{FC_R} as follows:

$$\begin{aligned} \Omega_{FC_R} &\stackrel{def}{=} \Delta_{X_{R_1}} \cup \dots \cup \Delta_{X_{R_n}} = \bigcup_{i=1}^n \Delta_{X_{R_i}} \\ &\text{where} \\ \Delta_{X_{R_i}} \cap \Delta_{X_{R_j}} &= \emptyset \quad \text{for } i, j \in [0, n] \mid i \neq j \\ \forall f \in \Delta_{X_{R_s}} \quad &f \text{ implements } f_{abstract} \text{ for } s \in [0, n] \end{aligned} \quad (11)$$

Since belong to different swarms, please note that any $\Delta_{X_{R_i}}$ and $\Delta_{X_{R_j}}$, for any $i \neq j$, can belong to the same provider but they cannot share the same concurrency limit.

Generally the schedulability condition for FC_R can be written as follows:

$$FC_R \text{ is schedulable} \Leftrightarrow \exists \Omega_{FC_R} \quad (12)$$

9.2 The Δ_{X_R} -Scheduler

Let R a resource owner, P the server-less provider and Δ_{X_R} a sub-swarm of a X_R , where k its concurrency limit.

Let $m \in \mathbb{N}$, a Δ_{X_R} -Scheduler, denoted as $S_{(\Delta_{X_R}, m)}$ represents a queuing system, implementing any scheduling discipline, equipped with m so-called *virtual function instance*, where $m \leq k$.

Its aim is to decide when and which function, belonging to Δ_{X_R} , must be performed on P .

The parameter m is also called *scheduler capacity*.

9.2.1 Virtual function instance

A *virtual function instance* represents a real function instances, clearly belonging to the server-less computing platform provider, which is *virtually* owned by $S_{(\Delta_{X_R}, m)}$.

Therefore, m represents the max number of server-less function instances usable simultaneously by $S_{(\Delta_{X_R}, m)}$.

9.2.2 Proprieties and constraints

According to our model, a Δ_{X_R} -scheduler capable to manage any function belonging to Δ_{X_R} , if exist, is *not* unique, although it is unique inside a peer node.

In order to achieve better performance in terms of network delay experienced by end users, fault tolerance and load balance, any peer nodes can hold a Δ_{X_R} -scheduler in order to manage incoming request sent by several users spread in different geographic regions.

However, despite there is no upper bound to the number of Δ_{X_R} -schedulers existing at the same time in our system, there is a limitation regarding the scheduler capacity of each existing scheduler.

Let's start summarizing all rules regarding Δ_{X_R} -schedulers:

1. All peer node of our system can hold a Δ_{X_R} -scheduler object.
2. Each node can hold only one instance of type Δ_{X_R} -scheduler.
3. Let $n \in \mathbb{N}$ such that $n \geq 1$, suppose that our system contains n peer nodes holding a Δ_{X_R} -scheduler.

To be more precise, let's say that a sequence $S_{(1,(\Delta_{X_R},m_1))}, \dots, S_{(i,(\Delta_{X_R},m_i))}$ exist at the same time in our system, where $S_{(i,(\Delta_{X_R},m_i))}$ represent the Δ_{X_R} -scheduler owned by i -th node having scheduler capacity equal to m_i .

Following constraint must be hold:

$$\sum_{i=1}^n m_i \leq k \quad (13)$$

where $k \in \mathbb{N}$, with $k > 0$, is the concurrency limit of the swarm X_R .

Remember that any sub-swarm Δ_{X_R} share the same concurrency limit of X_R . Therefore, equation 17 states that, the sum of all scheduler capacities which manage the functions belonging to Δ_{X_R} , must be less or equal to the max number of function instances executable at the same time on the server-less computing platform provider.

9.3 The FC_R -Scheduler

To support hybrid-scheduling, that is the ability to execute multiple concrete function implementations belonging to different providers or subjected to different concurrency limit, in order to select the most suitable concrete function implementation according to a given QoS, unfortunately only one “*scheduler*” is not enough.

We call *FC_R -Scheduler* a set of Δ_{X_R} -Schedulers where $\Delta_{X_R} \in \Omega_{FC_R}$. Is always required that $|FC_R| \geq 1$, that is, at least one scheduler must exist.

10 FC_R -Active Peer Node

According to our model, in order to effectively invoke all server-less concrete function belonging to a function choreography FC_R , is required that a peer node is “active”.

We said that a peer node A is *FC_R -active peer node*, or, simply, *active*, when it holds a Δ_{X_R} -scheduler for any sub-swarm in Ω_{FC_R} . Formally:

$$A \text{ is } FC_R\text{-active peer node} \Leftrightarrow \forall \Delta_{X_R} \in \Omega_{FC_R} \quad \exists S_{(\Delta_{X_R},m)} \text{ hold by } A \quad (14)$$

Multiple nodes can be active at the same time. Any node perform its scheduling decision independently.

11 Architecture overview

Our system design is based on a network of nodes, or *peers*, every of which has the *same functionality*; in fact, any of them is able to handle request submission, request scheduling and, potentially, request execution.

This is the reason according to which we can mark our proposal as a *P2P system*.

11.1 Overlay network

Our system's nodes are connected by an *overlay network*.

Definition 11.1 (Overlay network). An overlay network is a virtual network built on top of a physical network according to which each nodes can communicate with other if and only if they are connected by virtual links belonging to the virtual network.

Remark. A node may not be able to communicate directly with an arbitrary other node although they can communicate through physical network.

To be more precise, we have adopt a fully *centralized unstructured overlay network* because the *peer-resource index*, sometimes called *directory*, is centralized.

Please note that hybrid unstructured or fully decentralized solutions are technically possible, but guarantees about quality of service are very difficult.

11.1.1 Locality-awareness property

Locality-awareness is one of the essential characteristics of our system. In fact, if each peer is able to select his neighbours exploiting a suitable locality aware algorithm, is possible to decrease user experienced delays.

We have decided to adopt a multi-level based locality-aware neighbour selection called *intra-AS lowest delay clustering algorithm* (ASLDC).

When ASLDC algorithm is used, each peer chooses nearby peers only from those within the same AS; then it ranks its neighbours in terms of transmission latency, preferring to establish the connection with the node with the shortest latency to itself.

TODO [3]

11.2 Fault Tolerance

Like in many other P2P implementation, there are two ways of detecting failures in our proposed solution:

1. If a node tries to communicate with a neighbour and fails.
2. Since all nodes send to all his neighbour nodes so-called “heartbeat” messages, that is messages sent at fixed time intervals to indicate that the sender is alive, is possible to detect a failure by not receiving aforementioned periodic update messages after a long time.

11.2.1 Availability

Let $n \in \mathbb{N}$, R a resource owner and FC_R his function choreography, since multiple FC_R -Scheduler object can exist in n different nodes, our proposed solution can guarantee an high degree of availability.

In fact, when a node holding an FC_R -Scheduler fails, all end users requests related to FC_R can be routed to any other node holding a FC_R -Scheduler.

11.2.2 Replication

Although multiple FC_R -Scheduler object can exist, it's not mean that FC_R -Scheduler object are replicated, because every of them manage different virtual function instances.

Except the peer-resource index, any form of replication is performed by our system.

Algorithm 1 An algorithm with caption

$y \leftarrow 1$

and Δ_{X_R} -Scheduler

The locality awareness of the overlay network is used in scheduling jobs, described in Section 3.2

physical network but with added properties such as fault tolerance and flexibility.

The first problem is scheduling. Since there is no central scheduler it is difficult

Hybrid unstructured overlay

11.3 FC_R -Request

object with only one limitation:

Suppose that globally there are a set of schedulers $S_{1,(R_X,m_1)}, \dots, S_{p,(R_X,m_p)}$, where $p \in \mathbb{N}$ with $p \geq 1$

To be more precise, when a function x_j must to be execute, let s the current number of busy virtual instances, one of the following events may occur:

1. if $s < m$, the scheduler invoke directly the function x_j on the provider.
2. if $s = m$, the scheduler delay the execution of the function x_j on the provider according to implemented scheduling discipline.

Let R a resource owner and R_x its function choreography made up of $R_{x_1}, R_{x_2}, \dots, R_{x_n}$ unique server-less functions; it is said that a peer node P is **responsible** for R_x when it contains a sequence of schedulers $S_{R_1}, S_{R_2}, \dots, S_{R_k}$ with $k \leq n$, belonging to R , capable to invoke all server-less function belonging to R_x .

It is said that a

Depending on the definition of the function choreography provided by R and the unique characteristics of back-end server-less providers which execute all serverless functions R_{x_n} of

It is said that a scheduler S is capable to invoke a server-less function when , a scheduler S can invoke multiple

When a peer A , placed “*at the edge*” of the network, receives a new request of invocation for X by an end user, it performs following task in that order:

1. If it responsible It check for it is an already an *active peer* to manage

has found the tracker for a file F , the tracker returns a subset of all the nodes currently involved in downloading F .

Replication and Fault Tolerance. There are two ways of detecting failures in CAN, the first if a node tries to communicate with a neighbor and fails, it takes over that neighbor’s zone. The second way of detecting a failure is by not receiving the periodic update message after a long time. In the second case, the failure would probably be detected by all the neighbors, and all of them would try to take over the zone of the failed node, to resolve this, all nodes send to all other neighbors the size of their zone, and the node with the smallest zone takes over.

$$E[T] = \sum_{i=0}^n E[S_i] + E[T_{Q_i}] \quad (15)$$

12 Queuing system of a Δ_{X_R} -scheduler

Let $k \in \mathbb{N}$ such that $k \geq 0$. Moreover, suppose that R represents a resource owner, P a server-less provider and Δ_{X_R} a sub-swarm of the swarm X_R where k its concurrency limit. Finally, let $S_{(\Delta_{X_R}, m)}$ a Δ_{X_R} -scheduler.

Obliviously, since there are only m available virtual function instances, if there are more than m server-less functions waiting to be execute, a choice has to be made about which server-less function has to run next to ensuring QoS guarantees for latency critical applications.

Because we expect to execute server-less functions with different response-time requirements, which may have different scheduling needs, i have designed the Δ_{X_R} -scheduler as a queuing system implementing a *multilevel queue scheduling algorithm* which partitions the ready queue into several separate queues.

According to our solution, each queue has its own scheduling algorithm and any server-less function is permanently assigned to one queue according to his class.

In addition, we have adopt *round-robin* (RR) scheduling algorithm to perform scheduling activity among the queues.

12.1 Server-less function preemption

A very important consideration regards server-less function preemption.

In our context, due to FaaS paradigm, which hides the complexity of servers where our functions will be executed, the ability to preempt functions is not naturally available.

For that reason, only non-preemptive scheduling algorithms, according to which once a server-less function starts running it cannot be preempted, even if a higher priority server-less function comes along, can be adopted.

Since many scheduling algorithms require job preemption to run optimally, this situation can lead to a suboptimal resource management.

12.2 Queuing system design

A Δ_{X_R} -scheduler is made up of three queues:

- One queue implements a *Non-Preemptive Least-Slack-Time-First* (LST) scheduling algorithm. We will refer to that queue using Q_{LST} notation.

That algorithm implements a *dynamic priority scheduling approach* where priorities are assigned to server-less functions based on their *slacks*.

At any time t , the *slack* of a job with deadline at d is equal to $d - t - s$, where s is the time required to complete the job.

Any job having strict latency requirements must to be assign to this queue.

- Another queue implements, instead, a *Non-Preemptive Shortest-Job-First* (SJF) scheduling algorithm. That queue is denoted with Q_{SJF} .

That policy assigns priorities to jobs based on their size: the smaller the size, the higher the priority.

This queue is reserved for any function choreography that tolerates high latency. However, according to queueing theory, since SJF performs very bad when service time distribution is heavy-tailed; in fact

$$E[T_Q(x)]^{SJF} = \frac{\rho E[S^2]}{2E[S]} \cdot \frac{1}{(1 - \rho x)^2} \quad (16)$$

This queue is not suitable for very, very large server-less function. The variance in the job size distribution must be low.

- Finally, the third queue exploits the the simplest scheduling algorithm, that is the *First-Come-First-Served* (FCFS) policy, and it is denoted using Q_{FCFS} notation.

This queue is used for any heavy-tailed server-less function which can tolerate high latency.

12.3 The virtual function instances allocation problem

Let $n \in \mathbb{N}$ such that $n \geq 1$ and suppose that, at a certain time, a sequence including n unique Δ_{X_R} -schedulers exist in our system.

Formally, that sequence is denoted with $S_{(1,(\Delta_{X_R},m_1))}, \dots, S_{(i,(\Delta_{X_R},m_i))}, \dots, S_{(n,(\Delta_{X_R},m_n))}$, where $S_{(i,(\Delta_{X_R},m_i))}$ represents the Δ_{X_R} -scheduler owned by i -th node having scheduler capacity equal to m_i .

We have already established that the following constraint must be hold:

$$\sum_{i=1}^n m_i \leq k \quad (17)$$

where $k \in \mathbb{N}$, with $k > 0$, is the concurrency limit of the swarm X_R .

12.3.1 Reactive scaling policy

Every node of the system monitors all queues belong to its Δ_{X_R} -Scheduler and adjusts the number of reserved virtual function instances to preserve quality-of-service guarantees. This activity is called *reactive scaling policy*.

Although a $S_{(\Delta_{X_R},m)}$ owns m virtual function instances, they are not equally shared between aforementioned run queue.

First of all, let be v_{LST} , v_{SJF} and v_{FCFS} the numbers of the virtual function instances bounded to the Q_{LST} , Q_{SJF} and Q_{FCFS} run-queue.

At any time, following constrains must be hold:

$$v_{LST} \stackrel{def}{=} \frac{N \cdot E[S]}{\sum slack} \quad (18)$$

$$v_{SJF} \stackrel{def}{=} \frac{len(Q_{SJF})}{\alpha} \quad (19)$$

$$v_{FCFS} \stackrel{def}{=} \frac{len(Q_{FCFS})}{\alpha} \quad (20)$$

$$m = v_{LST} + v_{SJF} + v_{FCFS} \quad (21)$$

1. If the current number of virtual function instances is not enough to avoid deadline misses by latency constrained server-less function inside Q_{LST} , will be add to $VCPU_{LST}$ a number of virtual processor equal to

$$v_{LST} < \frac{N \cdot E[S]}{\sum slack} \quad (22)$$

This form of aging prevents starvation.

Finally, suppose that a node peer A holds a $S_{(\Delta_{X_R}, m_A)}$, which represents the Δ_{X_R} -Scheduler object.

1. Once a request is received over HTTP
2. Ask to the coordinator for

$$m = \frac{\sum slack}{E[S] \cdot m_{current}} + \frac{N}{\alpha} \quad (23)$$

$$\lambda_{FC_R} \geq \lambda_{threshold} \quad (24)$$

$$RTT_{(A,B)} > RTT_{(A,C)} + \quad (25)$$

$$f_{abstract} = next \quad (26)$$

$$f_{abstract} = \min(E[f_{\Delta_{X_{R_i}}}] + E[T_{Q_{LST}}]F \quad (27)$$

12.3.2 Server-less function scheduling for QoS “Minimum Response Time”

Let $n \in \mathbb{N}$, such that $n \geq 1$, supposing to have an FC_R -active peer node and let Ω_{FC_R} the set of all sub-swarms containing at least one function implementing any $f_{abstract} \in \mathbf{F}_{abstract}$, where $|\Omega_{FC_R}| = n$.

Finally, suppose that $f_{abstract}$ is the next abstract function that have to be executed according to the control-flow logic described by FC_R .

1. For $1 \leq i \leq n$, let $\Delta_{(X_{R_i}, f_{abstract})} \subseteq \Delta_{X_{R_i}}$ containing all concrete functions implementing $f_{abstract}$.
2. For $1 \leq i \leq n$, select $\mathbf{f}_i \in \Delta_{(X_{R_i}, f_{abstract})}$ functions having the minimum response time $E[T_{\mathbf{f}_i}]$.
3. Found $i \in [0, m]$ such that:

$$E[T_{\mathbf{f}_i}] + E[T_{Q_i}]^{\mathbf{LST}} < E[T_{\mathbf{f}_j}] + E[T_{Q_j}]^{\mathbf{LST}} \quad \text{for } j \in \mathbb{N}, 1 \leq j \leq n, i \neq j \quad (28)$$

4. Select the i -th Δ_{X_R} -Schedulers belonging to FC_R -Scheduler of the node and add \mathbf{f}_i into the Q_{LST} run-queue.